

# Diamonds

Francisco Arrieta, Emily Schmidt and Lucia Camenisch

2022-12-17

## Contents

|   |           |
|---|-----------|
| <b>Data Exploration</b>                                     | <b>2</b>  |
| Dimension Summary . . . . .                                 | 2         |
| Missing Values . . . . .                                    | 3         |
| Variable Visualisation . . . . .                            | 7         |
| <b>Variable Prediction and Model Performance Evaluation</b> | <b>8</b>  |
| Linear Regression . . . . .                                 | 8         |
| $k$ -NN . . . . .   | 24        |
| Regression Tree . . . . .                                   | 26        |
| NeuralNetworks . . . . .                                    | 33        |
| Ensembles . . . . .   | 40        |
| <b>Model Performance Summary</b>                            | <b>42</b> |
| <b>Conclusions</b>  | <b>43</b> |

### ##### General Use #####

```
library(car)           #for statistic functions
library(DataExplorer)  #for graphing missing value percentages
library(data.table)    #for reading data.tables
library(dplyr)         #for data manipulation
library(fastDummies)   #for creating dummies
library(e1071)         #for skewness
library(ellipse)       #for mapping correlation
library(GGally)        #for making graphs
library(ggplot2)       #for making graphs
library(ggpubr)        #for plot alignment
library(gridExtra)
library(kableExtra)    #for more elaborate tables
library(knitr)
library(tidyr)         #for changing the shape and hierarchy of a data set
library(naniar)        #for missing values
library(RColorBrewer)  #for graph colors
library(rattle)        #Graphical Data Interface
```

### ##### For Predictions #####

```
library(caret)         #for preProcess() and accuracy()
library(forecast)      # for accuracy() measures
library(FNN)           #for finding k nearest neighbor
library(gbm)           #for boosting
```

```

library(ipred)           #for bagging
library(vip)             #for variable importance
library(randomForest)    #for randomForest
library(rpart)           #for regression trees
library(rpart.plot)      #for plot trees
library(keras)           #front-end library for neural networks
library(magrittr)
library(tensorflow)      #backend python library for neural network

##### Personalized Functions #####
source("VIF.R")          #for calculating VIF (KEEP/DELETE?????????)
source("ProcStep.R")     #for variable selection (forw, backw, setpw)
source("GlobalCrit.R")   #for variable selection (exhaustive search)

##### Additonal Actions #####
options(scipen = 999)    #for removing scientific notation
tf$constant("Hello Tensorflow!") #for initializing tensorflow environment

## tf.Tensor(b'Hello Tensorflow!', shape=(), dtype=string)

```

## Data Exploration

```

diamonds <- fread("diamonds.csv", sep=",", header = T) # Load your data, diamonds.csv

diamonds$V1 <- NULL # Remove column 'V1' as it is similar to an ID variable - no additional meaning der

# Rename columns for more precise names
colnames(diamonds)[5] <- "depth_ratio" # depth to depth_ratio
colnames(diamonds)[8] <- "length" # x to length
colnames(diamonds)[9] <- "width" # y to width
colnames(diamonds)[10] <- "depth" # z to depth

```

## Dimension Summary

```

dim(diamonds) # Dimensions of data

## [1] 53940    10

summary(diamonds) # Produce result summaries of all variables

##      carat      cut      color      clarity
## Min.   :0.2000 Length:53940 Length:53940 Length:53940
## 1st Qu.:0.4000 Class :character Class :character Class :character
## Median :0.7000 Mode  :character Mode  :character Mode  :character
## Mean   :0.7979
## 3rd Qu.:1.0400
## Max.   :5.0100
## depth_ratio    table    price    length
## Min.   :43.00 Min.   :43.00 Min.   : 326 Min.   : 0.000
## 1st Qu.:61.00 1st Qu.:56.00 1st Qu.: 950 1st Qu.: 4.710
## Median :61.80 Median :57.00 Median : 2401 Median : 5.700
## Mean   :61.75 Mean   :57.46 Mean   : 3933 Mean   : 5.731

```

```
## 3rd Qu.:62.50 3rd Qu.:59.00 3rd Qu.: 5324 3rd Qu.: 6.540
## Max. :79.00 Max. :95.00 Max. :18823 Max. :10.740
## width depth
## Min. : 0.000 Min. : 0.000
## 1st Qu.: 4.720 1st Qu.: 2.910
## Median : 5.710 Median : 3.530
## Mean : 5.735 Mean : 3.539
## 3rd Qu.: 6.540 3rd Qu.: 4.040
## Max. :58.900 Max. :31.800
```

```
str(diamonds) # Type of variables
```

```
## Classes 'data.table' and 'data.frame': 53940 obs. of 10 variables:
## $ carat : num 0.23 0.21 0.23 0.29 0.31 0.24 0.24 0.26 0.22 0.23 ...
## $ cut : chr "Ideal" "Premium" "Good" "Premium" ...
## $ color : chr "E" "E" "E" "I" ...
## $ clarity : chr "SI2" "SI1" "VS1" "VS2" ...
## $ depth_ratio: num 61.5 59.8 56.9 62.4 63.3 62.8 62.3 61.9 65.1 59.4 ...
## $ table : num 55 61 65 58 58 57 57 55 61 61 ...
## $ price : int 326 326 327 334 335 336 336 337 337 338 ...
## $ length : num 3.95 3.89 4.05 4.2 4.34 3.94 3.95 4.07 3.87 4 ...
## $ width : num 3.98 3.84 4.07 4.23 4.35 3.96 3.98 4.11 3.78 4.05 ...
## $ depth : num 2.43 2.31 2.31 2.63 2.75 2.48 2.47 2.53 2.49 2.39 ...
## - attr(*, ".internal.selfref")=<externalptr>
```

```
# Number of unique values in each variable
sapply(diamonds, function(x) length(unique(x)))
```

```
## carat cut color clarity depth_ratio table
## 273 5 7 8 184 127
## price length width depth
## 11602 554 552 375
```

## Missing Values

```
# Missing values analysis
gg_miss_var(diamonds) + ggtitle("Missing values")
```

```
##### carat no problems #####
```

```
##### cut #####
unique(diamonds$cut) # Review unique values for cut
```

```
## [1] "Ideal" "Premium" "Good" "Very Good" "Fair"
```

```
# Factor the cut to five level
diamonds$cut <- as.factor(diamonds$cut)
```

```
# Ordered from worst to best
diamonds$cut <- ordered(diamonds$cut, levels = c("Fair", "Good", "Very Good", "Premium", "Ideal"))
```

```
##### color #####
# Review unique values for color
unique(diamonds$color)
```

```
## [1] "E" "I" "J" "H" "F" "G" "D"
# Factor the color to seven levels
diamonds$color <- as.factor(diamonds$color)

# Ordered from worst to best
diamonds$color <- ordered(diamonds$color, levels = c("J", "I", "H", "G", "F", "E", "D"))

##### clarity #####
# Review unique values for clarity
unique(diamonds$clarity)

## [1] "SI2" "SI1" "VS1" "VS2" "VVS2" "VVS1" "I1" "IF"
# Factor the clarity to eight levels
diamonds$clarity <- as.factor(diamonds$clarity)

# Ordered from worst to best
diamonds$clarity <- ordered(diamonds$clarity, levels = c("I1", "SI2", "SI1", "VS2", "VS1", "VVS2", "VVS1", "IF"))

##### table no problems#####

##### price no problems#####

##### Other Checks #####

# Remove values of 0 for for dimensions which includes zeros in length and width
nrow(diamonds[depth %in% 0,]) # Remove 20 rows due to depth = 0.0

## [1] 20
diamonds <- diamonds[depth > 0, ] # Include only values with depth greater than zero

# Create formula to check the absolute value of length to width, comparison
diamonds[, subtraction := abs(length - width)]
nrow(diamonds[subtraction>10,]) # Remove 2 rows due their extreme subtraction value (~59 and ~26)

## [1] 2
diamonds <- diamonds[subtraction <= 10, ] # Include only values with subtraction less than ten

# Check if the Depth_Ratio value corresponds to formula indicated in the description
diamonds[, depth_check := round(100*(2*depth)/((length + width)), 1)]
diamonds[, diff := abs(depth_check-depth_ratio)]

# Create histogram to look at the differences between how much price is off between a calculated value
hist(diamonds[diff >= 0.2 & diff < 1, diff], breaks =50, col = "#D8B365", border = "#D8B365", main = "TI")

# Threshold set to 0.3 due to ... (report)
nrow(diamonds[diff > 0.3,]) # We remove 268 rows

## [1] 253
diamonds <- diamonds[diff <= 0.3,]
```

```
# Removed created columns needed to clean the data
```

```
diamonds[, subtraction := NULL]
```

```
diamonds[, depth_check := NULL]
```

```
diamonds[, diff := NULL]
```

```
# Total rows remove: 275 observations
```

```
# Reorder data table to group like variable types
```

```
diamonds <- diamonds[, c(7, 2:4, 1, 8:10, 5:6)]
```

```
#Used ggpairs to create a scatterplot matrix between quantitative variables
```

```
ggpairs(diamonds[, c(1, 5:10)], title = "Scatterplot Matrix",
```

```
  proportions = "auto",
```

```
  columnLabels = c("Price", "Carat", "Length", "Width", "Depth", "Depth Ratio", "Table"),
```

```
  upper = list(continuous = wrap('cor', size = 3)),) + theme_light()
```

```
# Create plot that looks at carat and price
```

```
PC11 <- ggplot(aes(x = carat, y = price), data = diamonds) + geom_point(alpha = 0.5, size = 1, position = 'jitter') +  
  scale_color_brewer(type = 'div', guide = guide_legend(title = 'Clarity', reverse = T, override.aes = list(alpha = 0.5)))
```

```
# Create plot that looks at length and price
```

```
PC12 <- ggplot(aes(x = length, y = price), data = diamonds) + geom_point(alpha = 0.5, size = 1, position = 'jitter') +  
  scale_color_brewer(type = 'div', guide = guide_legend(title = 'Clarity', reverse = T, override.aes = list(alpha = 0.5)))
```

```
# Create plot that looks at width and price
```

```
PC13 <- ggplot(aes(x = width, y = price), data = diamonds) + geom_point(alpha = 0.5, size = 1, position = 'jitter') +  
  scale_color_brewer(type = 'div', guide = guide_legend(title = 'Clarity', reverse = T, override.aes = list(alpha = 0.5)))
```

```
# Create plot that looks at depth and price
```

```
PC14 <- ggplot(aes(x = depth, y = price), data = diamonds) + geom_point(alpha = 0.5, size = 1, position = 'jitter') +  
  scale_color_brewer(type = 'div', guide = guide_legend(title = 'Clarity', reverse = T, override.aes = list(alpha = 0.5)))
```

```
# Create plot that looks at depth_ratio and price
```

```
PC15 <- ggplot(aes(x = depth_ratio, y = price), data = diamonds) + geom_point(alpha = 0.5, size = 1, position = 'jitter') +
```

```
# Create plot that looks at table and price
```

```
PC16 <- ggplot(aes(x = table, y = price), data = diamonds) + geom_point(alpha = 0.5, size = 1, position = 'jitter') +
```

```
# Arrange ggplots into one frame
```

```
ggarrange(PC11, PC12, PC13, PC14, PC15, PC16,  
          ncol = 2, nrow = 3)
```

```
# Create plot that looks at carat and price
```

```
PCo1 <- ggplot(aes(x = carat, y = price), data = diamonds) + geom_point(alpha = 0.5, size = 1, position = 'jitter') +  
  scale_color_brewer(type = 'div', guide = guide_legend(title = 'Color', reverse = T, override.aes = list(alpha = 0.5)))
```

```
# Create plot that looks at length and price
```

```
PCo2 <- ggplot(aes(x = length, y = price), data = diamonds) + geom_point(alpha = 0.5, size = 1, position = 'jitter') +  
  scale_color_brewer(type = 'div', guide = guide_legend(title = 'Color', reverse = T, override.aes = list(alpha = 0.5)))
```

```
# Create plot that looks at width and price
```

```
PCo3 <- ggplot(aes(x = width, y = price), data = diamonds) + geom_point(alpha = 0.5, size = 1, position = 'jitter') +  
  scale_color_brewer(type = 'div', guide = guide_legend(title = 'Color', reverse = T, override.aes = list(alpha = 0.5)))
```

```

# Create plot that looks at depth and price
PCo4 <- ggplot(aes(x = depth, y = price), data = diamonds) + geom_point(alpha = 0.5, size = 1, position = 'jitter') +
  scale_color_brewer(type = 'div', guide = guide_legend(title = 'Color', reverse = T, override.aes = list(alpha = 0.5)))

# Create plot that looks at depth_ratio and price
PCo5 <- ggplot(aes(x = depth_ratio, y = price), data = diamonds) + geom_point(alpha = 0.5, size = 1, position = 'jitter') +
  scale_color_brewer(type = 'div', guide = guide_legend(title = 'Color', reverse = T, override.aes = list(alpha = 0.5)))

# Create plot that looks at table and price
PCo6 <- ggplot(aes(x = table, y = price), data = diamonds) + geom_point(alpha = 0.5, size = 1, position = 'jitter') +
  scale_color_brewer(type = 'div', guide = guide_legend(title = 'Color', reverse = T, override.aes = list(alpha = 0.5)))

# Arrange ggplots into one frame
ggarrange(PCo1, PCo2, PCo3, PCo4, PCo5, PCo6,
          ncol = 2, nrow = 3)

```

```

# Create plot that looks at carat and price
PCu1 <- ggplot(aes(x = carat, y = price), data = diamonds) + geom_point(alpha = 0.5, size = 1, position = 'jitter') +
  scale_color_brewer(type = 'div', guide = guide_legend(title = 'Cut', reverse = T, override.aes = list(alpha = 0.5)))

# Create plot that looks at length and price
PCu2 <- ggplot(aes(x = length, y = price), data = diamonds) + geom_point(alpha = 0.5, size = 1, position = 'jitter') +
  scale_color_brewer(type = 'div', guide = guide_legend(title = 'Cut', reverse = T, override.aes = list(alpha = 0.5)))

# Create plot that looks at width and price
PCu3 <- ggplot(aes(x = width, y = price), data = diamonds) + geom_point(alpha = 0.5, size = 1, position = 'jitter') +
  scale_color_brewer(type = 'div', guide = guide_legend(title = 'Cut', reverse = T, override.aes = list(alpha = 0.5)))

# Create plot that looks at depth and price
PCu4 <- ggplot(aes(x = depth, y = price), data = diamonds) + geom_point(alpha = 0.5, size = 1, position = 'jitter') +
  scale_color_brewer(type = 'div', guide = guide_legend(title = 'Cut', reverse = T, override.aes = list(alpha = 0.5)))

# Create plot that looks at depth_ratio and price
PCu5 <- ggplot(aes(x = depth_ratio, y = price), data = diamonds) + geom_point(alpha = 0.5, size = 1, position = 'jitter') +
  scale_color_brewer(type = 'div', guide = guide_legend(title = 'Cut', reverse = T, override.aes = list(alpha = 0.5)))

# Create plot that looks at table and price
PCu6 <- ggplot(aes(x = table, y = price), data = diamonds) + geom_point(alpha = 0.5, size = 1, position = 'jitter') +
  scale_color_brewer(type = 'div', guide = guide_legend(title = 'Cut', reverse = T, override.aes = list(alpha = 0.5)))

# Arrange ggplots into one frame
ggarrange(PCu1, PCu2, PCu3, PCu4, PCu5, PCu6,
          ncol = 2, nrow = 3)

```

```

# Correlation between price and quantitative variables
price_correlation <- with(diamonds,
  data.frame(cor_length_price = cor(length, price), cor_width_price = cor(width, price), cor_depth_price = cor(depth, price)))

# Transpose data and put into kable format
transpose <- t(sort(round(price_correlation, 4), decreasing = FALSE))
kable_corr <- kable(transpose) %>% kable_classic()
kable_corr

```

|                       |         |
|-----------------------|---------|
| cor_depth_ratio_price | -0.0104 |
| cor_table_price2      | 0.1273  |
| cor_depth_price       | 0.8829  |
| cor_length_price      | 0.8876  |
| cor_width_price       | 0.8892  |
| cor_carat_price3      | 0.9218  |

## Variable Visualisation

```
ggplot(gather(data = diamonds[, c(1, 5:10)]), aes(value)) +
  geom_histogram(aes(y = after_stat(density)),
    bins = 10,
    color = "white",
    fill = "#D8B365") + # Creates bin sizing and sets the lines as white
  geom_density(alpha = .2, fill = "#D8B365") +
  facet_wrap(~ key, scales = "free") + # Converting the graphs into panels
  ggtitle("Quantitative Variable Analysis") + # Title name
  ylab("Count") + xlab("Value") + # Label names
  theme_classic() # A classic theme, with x and y axis lines and no grid lines
```

```
# Create heatmap to show variable correlation
# Round the correlation coefficient to two decimal places
cormat <- round(cor(diamonds[, c(1, 5:10)]), 2)

# Use correlation between variables as distance
reorder_cormat <- function(cormat){
  dd <- as.dist((1-cormat)/2)
  hc <- hclust(dd)
  cormat <- cormat[hc$order, hc$order]
  return(cormat)
}

# Reorder the correlation matrix
cormat <- reorder_cormat(cormat)

# Keeping only upper triangular matrix
# upper_tri returns TRUE/FALSE for each coordinate (TRUE -> part of upper triangle)
# multiplying will thus keep the upper triangle values and set the others to 0
cormat <- cormat*upper.tri(cormat, diag = TRUE)
# Values of the lower triangle (0) are replaced by NA
cormat[cormat == 0] <- NA

# Melt the correlation matrix
cormat <- reshape2::melt(cormat, na.rm = TRUE)

# Create a ggheatmap with multiple characteristics
ggplot(cormat, aes(Var2, Var1, fill = value)) +
  geom_tile(color = "white") +
  scale_fill_gradient2(low = "#D8B365", high = "#15DDD8", mid = "white",
    midpoint = 0, limit = c(-1,1), space = "Lab", name="Pearson\nCorrelation") +
  ggtitle("Correlation Heatmap") + # Title name
  theme_minimal() + # Minimal theme, keeps in the lines
```

```
theme(axis.text.x = element_text(angle = 45, vjust = 1, size = 12, hjust = 1)) +
coord_fixed() +
geom_text(aes(Var2, Var1, label = value), color = "black", size = 2)
```

```
rm(cormat, reorder_cormat)
```

```
#plot Correlation ellipses
plotcorr(cor(diamonds[, -c(2:4)]), col = "#D8B365",
          main = "Pearson correlation ellipses for numerical variables")
```

```
# set seed for reproducing the partition
set.seed(111)

# generating training set index
train.index <- sample(c(1:nrow(diamonds)), 0.5*nrow(diamonds))

# generating validation set index taken from the complementary of training set
valid.index <- sample(setdiff(c(1:nrow(diamonds)), train.index), 0.3*nrow(diamonds))

# defining test set index as complementary of (train.index + valid.index)
test.index <- as.numeric(setdiff(row.names(diamonds), union(train.index, valid.index)))

# creating data tables Train, Valid and Test using the indexes
Train <- diamonds[train.index, ]
Valid <- diamonds[valid.index, ]
Test <- diamonds[test.index, ]
```

## Variable Prediction and Model Performance Evaluation

### Linear Regression

```
Train_lr <- diamonds[train.index, ]
Valid_lr <- diamonds[valid.index, ]
Test_lr <- diamonds[ test.index, ]
```

```
VIF(y = diamonds$price, matx = diamonds[, -c(1)])
```

```
##
##              GVIF Df GVIF^(1/(2*Df))
## cut          2.45522 4      1.11882
## color        1.18398 6      1.01417
## clarity      1.36848 7      1.02266
## carat        25.91780 1      5.09095
## length      1091.42000 1     33.03670
## width        1143.44000 1     33.81480
## depth        2008.33000 1     44.81440
## depth_ratio  31.99350 1      5.65628
## table        1.80396 1      1.34311
##
## Mean: 165.105
```



```
VIF(y = diamonds$price, matx = diamonds[, -c(1, 6, 7, 8)])
```

```
##
##              GVIF Df GVIF^(1/(2*Df))
## cut          1.93382 4          1.08593
## color        1.17045 6          1.01320
## clarity      1.30388 7          1.01913
## carat        1.32381 1          1.15057
## depth_ratio  1.38914 1          1.17862
## table        1.79505 1          1.33980
##
## Mean: 1.98352
```

```
plotcorr(corr(diamonds[, -c(2:4, 6:8)]), col = "#D8B365",
         main = "Pearson correlation ellipses for numerical variables")
```

```
sapply(Train_lr[, c(1, 5:10)], skewness)
```

```
##      price      carat      length      width      depth depth_ratio
## 1.62597473 1.09675143 0.39846889 0.39220078 0.39210931 0.01824979
##      table
## 0.87288692
```

```
Train_lr$price <- log(Train_lr$price)
Train_lr$carat <- log(Train_lr$carat)
Train_lr$table <- log(Train_lr$table)
sapply(Train_lr[, c(1, 5:10)], skewness)
```

```
##      price      carat      length      width      depth depth_ratio
## 0.11305083 0.09668960 0.39846889 0.39220078 0.39210931 0.01824979
##      table
## 0.64541910
```

```
Valid_lr$price <- log(Valid_lr$price)
Valid_lr$carat <- log(Valid_lr$carat)
Valid_lr$table <- log(Valid_lr$table)
```

```
ggplot(gather(data = Train_lr[, c(1, 5:10)]), aes(value)) +
  geom_histogram(aes(y = after_stat(density)),
                 color = "white",
                 fill = "#D8B365") + # Creates bin sizing with colors
  geom_density(alpha = .2, fill = "#D8B365") +
  facet_wrap(~ key, scales = "free") + # Converting the graphs into panels
  ggtitle("Histograms of numerical variables") + # Title name
  ylab("Count") + xlab("Value") + # Label names
  theme_classic() # Theme with x and y axis lines and no grid lines
```

```
norm.values <- preProcess(Train_lr[, c(1, 5:10)], method=c("center", "scale"))
```

```
Train_lr[, c(1, 5:10)] <- predict(norm.values, Train_lr[, c(1, 5:10)])
Valid_lr[, c(1, 5:10)] <- predict(norm.values, Valid_lr[, c(1, 5:10)])
Test_lr[, c(1, 5:10)] <- predict(norm.values, Test_lr[, c(1, 5:10)])
```

```
LM_complete = lm(price ~ . , data = Train_lr)
summary(LM_complete)
```

```
##
## Call:
## lm(formula = price ~ . , data = Train_lr)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1.04059 -0.08291  0.00028  0.08156  1.42487
##
## Coefficients:
##              Estimate Std. Error t value      Pr(>|t|)
## (Intercept) -0.0747373   0.0016289  -45.882 < 0.0000000000000002 ***
## cut.L         0.1168565   0.0037189   31.422 < 0.0000000000000002 ***
## cut.Q        -0.0342102   0.0030248  -11.310 < 0.0000000000000002 ***
## cut.C         0.0174498   0.0027023    6.457  0.0000000001084 ***
## cut^4         0.0003697   0.0020867    0.177    0.85937
## color.L       0.4366728   0.0028463  153.417 < 0.0000000000000002 ***
## color.Q      -0.0974458   0.0025856  -37.687 < 0.0000000000000002 ***
## color.C       0.0126123   0.0024180    5.216  0.0000001841568 ***
## color^4       0.0134193   0.0022210    6.042  0.0000000015434 ***
## color^5       0.0001336   0.0021044    0.063    0.94940
## color^6       0.0026592   0.0019108    1.392    0.16404
## clarity.L     0.9080235   0.0050105  181.224 < 0.0000000000000002 ***
## clarity.Q    -0.2486235   0.0046594  -53.360 < 0.0000000000000002 ***
## clarity.C     0.1370069   0.0039810   34.415 < 0.0000000000000002 ***
## clarity^4    -0.0685371   0.0031777  -21.568 < 0.0000000000000002 ***
## clarity^5     0.0290689   0.0025896   11.225 < 0.0000000000000002 ***
## clarity^6    -0.0022942   0.0022444   -1.022    0.30670
## clarity^7     0.0312779   0.0019833   15.771 < 0.0000000000000002 ***
## carat        0.9997686   0.0082204  121.621 < 0.0000000000000002 ***
## length       0.1775342   0.0266043    6.673  0.0000000000255 ***
## width       -0.0200231   0.0270778   -0.739    0.45963
## depth       -0.0702733   0.0359774   -1.953    0.05080 .
## depth_ratio  0.0117907   0.0045226    2.607    0.00914 **
## table        0.0011885   0.0010849    1.096    0.27330
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.131 on 26808 degrees of freedom
## Multiple R-squared:  0.9829, Adjusted R-squared:  0.9828
## F-statistic: 6.685e+04 on 23 and 26808 DF,  p-value: < 0.0000000000000002
```

```
LM_forward_complete = step(LM_complete, direction = "forward")
```

```
## Start: AIC=-109064.2
## price ~ cut + color + clarity + carat + length + width + depth +
##      depth_ratio + table
summary(LM_forward_complete)    # no selection
```

```
##
## Call:
## lm(formula = price ~ cut + color + clarity + carat + length +
```

```

##      width + depth + depth_ratio + table, data = Train_lr)
##
## Residuals:
##      Min        1Q      Median        3Q        Max
## -1.04059 -0.08291  0.00028  0.08156  1.42487
##
## Coefficients:
##              Estimate Std. Error t value      Pr(>|t|)
## (Intercept) -0.0747373  0.0016289 -45.882 < 0.0000000000000002 ***
## cut.L        0.1168565  0.0037189  31.422 < 0.0000000000000002 ***
## cut.Q       -0.0342102  0.0030248 -11.310 < 0.0000000000000002 ***
## cut.C        0.0174498  0.0027023   6.457  0.0000000001084 ***
## cut^4        0.0003697  0.0020867   0.177    0.85937
## color.L      0.4366728  0.0028463 153.417 < 0.0000000000000002 ***
## color.Q     -0.0974458  0.0025856 -37.687 < 0.0000000000000002 ***
## color.C      0.0126123  0.0024180   5.216  0.0000001841568 ***
## color^4      0.0134193  0.0022210   6.042  0.0000000015434 ***
## color^5      0.0001336  0.0021044   0.063    0.94940
## color^6      0.0026592  0.0019108   1.392    0.16404
## clarity.L    0.9080235  0.0050105 181.224 < 0.0000000000000002 ***
## clarity.Q   -0.2486235  0.0046594 -53.360 < 0.0000000000000002 ***
## clarity.C    0.1370069  0.0039810  34.415 < 0.0000000000000002 ***
## clarity^4   -0.0685371  0.0031777 -21.568 < 0.0000000000000002 ***
## clarity^5    0.0290689  0.0025896  11.225 < 0.0000000000000002 ***
## clarity^6   -0.0022942  0.0022444  -1.022    0.30670
## clarity^7    0.0312779  0.0019833  15.771 < 0.0000000000000002 ***
## carat       0.9997686  0.0082204 121.621 < 0.0000000000000002 ***
## length      0.1775342  0.0266043   6.673  0.0000000000255 ***
## width      -0.0200231  0.0270778  -0.739    0.45963
## depth      -0.0702733  0.0359774  -1.953    0.05080 .
## depth_ratio 0.0117907  0.0045226   2.607    0.00914 **
## table       0.0011885  0.0010849   1.096    0.27330
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.131 on 26808 degrees of freedom
## Multiple R-squared:  0.9829, Adjusted R-squared:  0.9828
## F-statistic: 6.685e+04 on 23 and 26808 DF,  p-value: < 0.00000000000000022
LM_backward_complete = step(LM_complete, direction = "backward")

## Start:  AIC=-109064.2
## price ~ cut + color + clarity + carat + length + width + depth +
##      depth_ratio + table
##
##              Df Sum of Sq    RSS    AIC
## - width        1      0.01  459.84 -109066
## - table         1      0.02  459.85 -109065
## <none>                      459.83 -109064
## - depth        1      0.07  459.89 -109062
## - depth_ratio  1      0.12  459.95 -109059
## - length       1      0.76  460.59 -109022
## - cut          4     19.38  479.21 -107965
## - carat        1    253.72  713.54  -97277
## - color        6    428.53  888.36  -91407

```

```

## - clarity      7      874.13 1333.96 -80501
##
## Step: AIC=-109065.7
## price ~ cut + color + clarity + carat + length + depth + depth_ratio +
##      table
##
##           Df Sum of Sq      RSS      AIC
## - table      1         0.02  459.86 -109066
## <none>                        459.84 -109066
## - depth      1         0.18  460.02 -109057
## - depth_ratio 1         0.30  460.14 -109050
## - length     1         0.75  460.59 -109024
## - cut        4        19.38  479.22 -107966
## - carat      1       256.52  716.35 -97173
## - color      6       428.96  888.80 -91395
## - clarity    7       877.59 1337.43 -80433
##
## Step: AIC=-109066.4
## price ~ cut + color + clarity + carat + length + depth + depth_ratio
##
##           Df Sum of Sq      RSS      AIC
## <none>                        459.86 -109066
## - depth      1         0.18  460.04 -109058
## - depth_ratio 1         0.29  460.15 -109052
## - length     1         0.75  460.61 -109025
## - cut        4        24.58  484.44 -107677
## - carat      1       261.25  721.11 -96998
## - color      6       429.14  889.00 -91391
## - clarity    7       877.84 1337.70 -80430
summary(LM_backward_complete)  # like LM_CpAIC_complete

##
## Call:
## lm(formula = price ~ cut + color + clarity + carat + length +
##      depth + depth_ratio, data = Train_lr)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1.04014 -0.08295  0.00025  0.08161  1.42593
##
## Coefficients:
##              Estimate Std. Error t value      Pr(>|t|)
## (Intercept) -0.0741493   0.0015667  -47.329 < 0.0000000000000002 ***
## cut.L        0.1150923   0.0034602   33.262 < 0.0000000000000002 ***
## cut.Q       -0.0339481   0.0029582  -11.476 < 0.0000000000000002 ***
## cut.C        0.0163231   0.0025452    6.413   0.0000000001447 ***
## cut^4       -0.0001739   0.0020487   -0.085     0.93235
## color.L      0.4366337   0.0028441  153.523 < 0.0000000000000002 ***
## color.Q     -0.0974187   0.0025833  -37.711 < 0.0000000000000002 ***
## color.C      0.0125915   0.0024179    5.208   0.0000001927323 ***
## color^4      0.0133674   0.0022204    6.020   0.0000000017651 ***
## color^5      0.0001194   0.0021042    0.057     0.95475
## color^6      0.0026739   0.0019108    1.399     0.16171
## clarity.L    0.9075881   0.0049934  181.757 < 0.0000000000000002 ***

```

```
## clarity.Q    -0.2483814  0.0046523 -53.389 < 0.0000000000000002 ***
## clarity.C     0.1367692  0.0039749  34.408 < 0.0000000000000002 ***
## clarity^4    -0.0684754  0.0031760 -21.560 < 0.0000000000000002 ***
## clarity^5     0.0289678  0.0025881  11.193 < 0.0000000000000002 ***
## clarity^6    -0.0022952  0.0022444  -1.023      0.30649
## clarity^7     0.0313111  0.0019831  15.789 < 0.0000000000000002 ***
## carat         1.0002821  0.0081052 123.413 < 0.0000000000000002 ***
## length        0.1751101  0.0264602   6.618      0.0000000000371 ***
## depth        -0.0884193  0.0271461  -3.257      0.00113 **
## depth_ratio   0.0136166  0.0033272   4.093      0.0000427910418 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.131 on 26810 degrees of freedom
## Multiple R-squared:  0.9829, Adjusted R-squared:  0.9828
## F-statistic: 7.321e+04 on 21 and 26810 DF,  p-value: < 0.00000000000000022
LM_stepwise_complete = step(LM_complete, direction = "both")
```

```
## Start:  AIC=-109064.2
## price ~ cut + color + clarity + carat + length + width + depth +
##      depth_ratio + table
##
##           Df Sum of Sq    RSS    AIC
## - width      1      0.01  459.84 -109066
## - table       1      0.02  459.85 -109065
## <none>                459.83 -109064
## - depth       1      0.07  459.89 -109062
## - depth_ratio  1      0.12  459.95 -109059
## - length       1      0.76  460.59 -109022
## - cut          4     19.38  479.21 -107965
## - carat        1    253.72  713.54 -97277
## - color        6    428.53  888.36 -91407
## - clarity      7    874.13 1333.96 -80501
##
## Step:  AIC=-109065.7
## price ~ cut + color + clarity + carat + length + depth + depth_ratio +
##      table
##
##           Df Sum of Sq    RSS    AIC
## - table       1      0.02  459.86 -109066
## <none>                459.84 -109066
## + width       1      0.01  459.83 -109064
## - depth       1      0.18  460.02 -109057
## - depth_ratio  1      0.30  460.14 -109050
## - length       1      0.75  460.59 -109024
## - cut          4     19.38  479.22 -107966
## - carat        1    256.52  716.35 -97173
## - color        6    428.96  888.80 -91395
## - clarity      7    877.59 1337.43 -80433
##
## Step:  AIC=-109066.4
## price ~ cut + color + clarity + carat + length + depth + depth_ratio
##
##           Df Sum of Sq    RSS    AIC
```

```

## <none>                459.86 -109066
## + table                1         0.02 459.84 -109066
## + width                1         0.01 459.85 -109065
## - depth               1         0.18 460.04 -109058
## - depth_ratio        1         0.29 460.15 -109052
## - length             1         0.75 460.61 -109025
## - cut                 4        24.58 484.44 -107677
## - carat              1       261.25 721.11 -96998
## - color              6       429.14 889.00 -91391
## - clarity            7       877.84 1337.70 -80430

summary(LM_stepwise_complete)  # like LM_CpAIC_complete

##
## Call:
## lm(formula = price ~ cut + color + clarity + carat + length +
##      depth + depth_ratio, data = Train_lr)
##
## Residuals:
##      Min        1Q    Median        3Q        Max
## -1.04014 -0.08295  0.00025  0.08161  1.42593
##
## Coefficients:
##              Estimate Std. Error t value      Pr(>|t|)
## (Intercept) -0.0741493  0.0015667 -47.329 < 0.0000000000000002 ***
## cut.L         0.1150923  0.0034602  33.262 < 0.0000000000000002 ***
## cut.Q        -0.0339481  0.0029582 -11.476 < 0.0000000000000002 ***
## cut.C         0.0163231  0.0025452   6.413  0.0000000001447 ***
## cut^4        -0.0001739  0.0020487  -0.085    0.93235
## color.L       0.4366337  0.0028441 153.523 < 0.0000000000000002 ***
## color.Q      -0.0974187  0.0025833 -37.711 < 0.0000000000000002 ***
## color.C       0.0125915  0.0024179   5.208  0.0000001927323 ***
## color^4       0.0133674  0.0022204   6.020  0.0000000017651 ***
## color^5       0.0001194  0.0021042   0.057    0.95475
## color^6       0.0026739  0.0019108   1.399    0.16171
## clarity.L     0.9075881  0.0049934 181.757 < 0.0000000000000002 ***
## clarity.Q    -0.2483814  0.0046523 -53.389 < 0.0000000000000002 ***
## clarity.C     0.1367692  0.0039749  34.408 < 0.0000000000000002 ***
## clarity^4    -0.0684754  0.0031760 -21.560 < 0.0000000000000002 ***
## clarity^5     0.0289678  0.0025881  11.193 < 0.0000000000000002 ***
## clarity^6    -0.0022952  0.0022444  -1.023    0.30649
## clarity^7     0.0313111  0.0019831  15.789 < 0.0000000000000002 ***
## carat        1.0002821  0.0081052 123.413 < 0.0000000000000002 ***
## length       0.1751101  0.0264602   6.618  0.0000000000371 ***
## depth       -0.0884193  0.0271461  -3.257    0.00113 **
## depth_ratio  0.0136166  0.0033272   4.093  0.0000427910418 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.131 on 26810 degrees of freedom
## Multiple R-squared:  0.9829, Adjusted R-squared:  0.9828
## F-statistic: 7.321e+04 on 21 and 26810 DF,  p-value: < 0.0000000000000002

options(scipen = 999)
GlobalCrit(LM_complete)

```

```

##
## -----
## GLOBAL VARIABLE SELECTION PROCEDURE
##
## ( Data = Train_lr )
##
## A = cut
## B = color
## C = clarity
## D = carat
## E = length
## F = width
## G = depth
## H = depth_ratio
## I = table
##
## Models      | Cp              | AIC              |
## -----
## ABCDEF      | 2.84 (10) | 32937.47 (10) |
## ABCDEH      | 2.49 ( 9) | 32937.81 ( 9) |
## ABCDEFG     | - 0.63 ( 7) | 32940.94 ( 7) |
## ABCDEFH     | - 3.12 ( 5) | 32943.44 ( 5) |
## ABCDEGH     | - 6.12 ( 1) | 32946.43 ( 1) |
## ABCDEFGH    | - 4.80 ( 3) | 32945.11 ( 3) |
## ABCDEFGI    | 0.80 ( 8) | 32939.51 ( 8) |
## ABCDEFHI    | - 2.18 ( 6) | 32942.50 ( 6) |
## ABCDEGHI    | - 5.45 ( 2) | 32945.77 ( 2) |
## ABCDEFGHI   | - 3.00 ( 4) | 32944.32 ( 4) |
##
## -----
##
## -----
## GLOBAL VARIABLE SELECTION PROCEDURE
##
## ( Data = Train_lr )
##
## A = cut
## B = color
## C = clarity
## D = carat
## E = length
## F = width
## G = depth
## H = depth_ratio
## I = table
##
## Models      | Cp              | AIC              |
## -----
## ABCDEF      | 2.84 (10) | 32937.47 (10) |
## ABCDEH      | 2.49 ( 9) | 32937.81 ( 9) |
## ABCDEFG     | - 0.63 ( 7) | 32940.94 ( 7) |
## ABCDEFH     | - 3.12 ( 5) | 32943.44 ( 5) |
## ABCDEGH     | - 6.12 ( 1) | 32946.43 ( 1) |

```

```

##   ABCDEFGH   | -   4.80 ( 3) | 32945.11 ( 3) |
##   ABCDEFGI   |      0.80 ( 8) | 32939.51 ( 8) |
##   ABCDEFHI   | -   2.18 ( 6) | 32942.50 ( 6) |
##   ABCDEGHI   | -   5.45 ( 2) | 32945.77 ( 2) |
##   ABCDEFGHI  | -   3.00 ( 4) | 32944.32 ( 4) |
##
## -----
LM_CpAIC_complete = lm(price ~ ., data = Train_lr[, c(1:6, 8, 9)])
summary(LM_CpAIC_complete)

##
## Call:
## lm(formula = price ~ ., data = Train_lr[, c(1:6, 8, 9)])
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1.04014 -0.08295  0.00025  0.08161  1.42593
##
## Coefficients:
##              Estimate Std. Error t value      Pr(>|t|)
## (Intercept) -0.0741493   0.0015667  -47.329 < 0.0000000000000002 ***
## cut.L         0.1150923   0.0034602   33.262 < 0.0000000000000002 ***
## cut.Q        -0.0339481   0.0029582  -11.476 < 0.0000000000000002 ***
## cut.C         0.0163231   0.0025452    6.413   0.0000000001447 ***
## cut^4        -0.0001739   0.0020487   -0.085     0.93235
## color.L       0.4366337   0.0028441  153.523 < 0.0000000000000002 ***
## color.Q      -0.0974187   0.0025833  -37.711 < 0.0000000000000002 ***
## color.C       0.0125915   0.0024179    5.208   0.0000001927323 ***
## color^4       0.0133674   0.0022204    6.020   0.0000000017651 ***
## color^5       0.0001194   0.0021042    0.057     0.95475
## color^6       0.0026739   0.0019108    1.399     0.16171
## clarity.L     0.9075881   0.0049934  181.757 < 0.0000000000000002 ***
## clarity.Q    -0.2483814   0.0046523  -53.389 < 0.0000000000000002 ***
## clarity.C     0.1367692   0.0039749   34.408 < 0.0000000000000002 ***
## clarity^4    -0.0684754   0.0031760  -21.560 < 0.0000000000000002 ***
## clarity^5     0.0289678   0.0025881   11.193 < 0.0000000000000002 ***
## clarity^6    -0.0022952   0.0022444   -1.023     0.30649
## clarity^7     0.0313111   0.0019831   15.789 < 0.0000000000000002 ***
## carat        1.0002821   0.0081052  123.413 < 0.0000000000000002 ***
## length       0.1751101   0.0264602    6.618   0.0000000000371 ***
## depth       -0.0884193   0.0271461   -3.257     0.00113 **
## depth_ratio  0.0136166   0.0033272    4.093   0.0000427910418 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.131 on 26810 degrees of freedom
## Multiple R-squared:  0.9829, Adjusted R-squared:  0.9828
## F-statistic: 7.321e+04 on 21 and 26810 DF,  p-value: < 0.00000000000000022

Train_minus_corr <- Train_lr[, -c(6:8)]
LM_minus_corr = lm(price ~ ., data = Train_minus_corr)
summary(LM_minus_corr)

##

```



```

## Call:
## lm(formula = price ~ ., data = Train_minus_corr)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.98683 -0.08456 -0.00054  0.08206  1.42888
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -0.0733380  0.0016248 -45.136 < 0.0000000000000002 ***
## cut.L         0.1174782  0.0037111  31.656 < 0.0000000000000002 ***
## cut.Q        -0.0331237  0.0029650 -11.172 < 0.0000000000000002 ***
## cut.C         0.0132619  0.0025653   5.170  0.0000002361 ***
## cut^4        -0.0022453  0.0020453  -1.098    0.2723
## color.L       0.4322410  0.0028233 153.096 < 0.0000000000000002 ***
## color.Q      -0.0958030  0.0025851 -37.059 < 0.0000000000000002 ***
## color.C       0.0133140  0.0024238   5.493  0.0000000399 ***
## color^4       0.0125884  0.0022260   5.655  0.0000000157 ***
## color^5      -0.0001023  0.0021099  -0.048    0.9613
## color^6       0.0028366  0.0019160   1.481    0.1388
## clarity.L     0.9068901  0.0050015 181.324 < 0.0000000000000002 ***
## clarity.Q    -0.2447781  0.0046522 -52.615 < 0.0000000000000002 ***
## clarity.C     0.1346227  0.0039816  33.811 < 0.0000000000000002 ***
## clarity^4    -0.0686749  0.0031846 -21.564 < 0.0000000000000002 ***
## clarity^5     0.0288881  0.0025955  11.130 < 0.0000000000000002 ***
## clarity^6    -0.0024999  0.0022505  -1.111    0.2667
## clarity^7     0.0320730  0.0019874  16.138 < 0.0000000000000002 ***
## carat        1.0865804  0.0009254 1174.133 < 0.0000000000000002 ***
## depth_ratio -0.0016945  0.0009461  -1.791    0.0733 .
## table       -0.0001141  0.0010774  -0.106    0.9157
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.1313 on 26811 degrees of freedom
## Multiple R-squared:  0.9828, Adjusted R-squared:  0.9828
## F-statistic: 7.644e+04 on 20 and 26811 DF, p-value: < 0.00000000000000022
LM_backward_minus_corr = step(LM_minus_corr, direction = "backward")

## Start: AIC=-108920.2
## price ~ cut + color + clarity + carat + depth_ratio + table
##
##              Df Sum of Sq    RSS    AIC
## - table         1      0.0   462.4 -108922
## <none>              462.4 -108920
## - depth_ratio   1      0.1   462.5 -108919
## - cut           4     20.1   482.6 -107784
## - color         6    429.9   892.3  -91293
## - clarity       7    882.5  1344.9  -80288
## - carat         1  23776.4 24238.8  -2687
##
## Step: AIC=-108922.2
## price ~ cut + color + clarity + carat + depth_ratio
##
##              Df Sum of Sq    RSS    AIC

```

```

## <none>                462.4 -108922
## - depth_ratio 1      0.1  462.5 -108921
## - cut         4      26.3  488.7 -107445
## - color       6     430.0  892.4 -91293
## - clarity     7     883.1 1345.5 -80277
## - carat       1    24022.4 24484.8 -2418

summary(LM_backward_minus_corr)  # like LM_CpAIC_minus_corr

##
## Call:
## lm(formula = price ~ cut + color + clarity + carat + depth_ratio,
##     data = Train_minus_corr)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.98668 -0.08451 -0.00056  0.08207  1.42904
##
## Coefficients:
##              Estimate Std. Error t value      Pr(>|t|)
## (Intercept) -0.0733824  0.0015697  -46.750 < 0.0000000000000002 ***
## cut.L        0.1176217  0.0034546   34.048 < 0.0000000000000002 ***
## cut.Q       -0.0331066  0.0029605  -11.183 < 0.0000000000000002 ***
## cut.C        0.0133117  0.0025217    5.279   0.0000001310 ***
## cut^4       -0.0022101  0.0020180   -1.095     0.2734
## color.L      0.4322448  0.0028231  153.112 < 0.0000000000000002 ***
## color.Q     -0.0958088  0.0025845  -37.070 < 0.0000000000000002 ***
## color.C      0.0133162  0.0024237    5.494   0.0000000396 ***
## color^4      0.0125930  0.0022255    5.658   0.0000000154 ***
## color^5     -0.0001032  0.0021098   -0.049     0.9610
## color^6      0.0028362  0.0019159    1.480     0.1388
## clarity.L    0.9069041  0.0049997  181.393 < 0.0000000000000002 ***
## clarity.Q   -0.2447790  0.0046521  -52.616 < 0.0000000000000002 ***
## clarity.C    0.1346281  0.0039812   33.816 < 0.0000000000000002 ***
## clarity^4   -0.0686735  0.0031846  -21.565 < 0.0000000000000002 ***
## clarity^5    0.0288920  0.0025952   11.133 < 0.0000000000000002 ***
## clarity^6   -0.0024996  0.0022504   -1.111     0.2667
## clarity^7    0.0320720  0.0019873   16.138 < 0.0000000000000002 ***
## carat       1.0865705  0.0009207 1180.212 < 0.0000000000000002 ***
## depth_ratio -0.0016535  0.0008632   -1.916     0.0554 .
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.1313 on 26812 degrees of freedom
## Multiple R-squared:  0.9828, Adjusted R-squared:  0.9828
## F-statistic: 8.047e+04 on 19 and 26812 DF,  p-value: < 0.0000000000000002

LM_forward_minus_corr = step(LM_minus_corr, direction = "forward")

## Start:  AIC=-108920.2
## price ~ cut + color + clarity + carat + depth_ratio + table

summary(LM_forward_minus_corr)  # no selection

##
## Call:

```

```

## lm(formula = price ~ cut + color + clarity + carat + depth_ratio +
##      table, data = Train_minus_corr)
##
## Residuals:
##      Min        1Q    Median        3Q        Max
## -0.98683 -0.08456 -0.00054  0.08206  1.42888
##
## Coefficients:
##              Estimate Std. Error t value      Pr(>|t|)
## (Intercept) -0.0733380  0.0016248  -45.136 < 0.0000000000000002 ***
## cut.L         0.1174782  0.0037111   31.656 < 0.0000000000000002 ***
## cut.Q        -0.0331237  0.0029650  -11.172 < 0.0000000000000002 ***
## cut.C         0.0132619  0.0025653    5.170  0.0000002361 ***
## cut^4        -0.0022453  0.0020453   -1.098    0.2723
## color.L       0.4322410  0.0028233  153.096 < 0.0000000000000002 ***
## color.Q      -0.0958030  0.0025851  -37.059 < 0.0000000000000002 ***
## color.C       0.0133140  0.0024238    5.493  0.0000000399 ***
## color^4       0.0125884  0.0022260    5.655  0.0000000157 ***
## color^5      -0.0001023  0.0021099   -0.048    0.9613
## color^6       0.0028366  0.0019160    1.481    0.1388
## clarity.L     0.9068901  0.0050015  181.324 < 0.0000000000000002 ***
## clarity.Q    -0.2447781  0.0046522  -52.615 < 0.0000000000000002 ***
## clarity.C     0.1346227  0.0039816   33.811 < 0.0000000000000002 ***
## clarity^4    -0.0686749  0.0031846  -21.564 < 0.0000000000000002 ***
## clarity^5     0.0288881  0.0025955   11.130 < 0.0000000000000002 ***
## clarity^6    -0.0024999  0.0022505   -1.111    0.2667
## clarity^7     0.0320730  0.0019874   16.138 < 0.0000000000000002 ***
## carat        1.0865804  0.0009254 1174.133 < 0.0000000000000002 ***
## depth_ratio -0.0016945  0.0009461   -1.791    0.0733 .
## table        -0.0001141  0.0010774   -0.106    0.9157
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.1313 on 26811 degrees of freedom
## Multiple R-squared:  0.9828, Adjusted R-squared:  0.9828
## F-statistic: 7.644e+04 on 20 and 26811 DF, p-value: < 0.00000000000000022
LM_stepwise_minus_corr = step(LM_minus_corr, direction = "both")

## Start: AIC=-108920.2
## price ~ cut + color + clarity + carat + depth_ratio + table
##
##              Df Sum of Sq    RSS    AIC
## - table         1      0.0  462.4 -108922
## <none>              462.4 -108920
## - depth_ratio   1      0.1  462.5 -108919
## - cut           4     20.1  482.6 -107784
## - color         6    429.9  892.3  -91293
## - clarity       7    882.5 1344.9  -80288
## - carat         1  23776.4 24238.8  -2687
##
## Step: AIC=-108922.2
## price ~ cut + color + clarity + carat + depth_ratio
##
##              Df Sum of Sq    RSS    AIC

```

```

## <none>                                462.4 -108922
## - depth_ratio 1          0.1  462.5 -108921
## + table      1          0.0  462.4 -108920
## - cut        4         26.3  488.7 -107445
## - color      6        430.0  892.4 -91293
## - clarity    7        883.1 1345.5 -80277
## - carat      1       24022.4 24484.8 -2418

summary(LM_stepwise_minus_corr)  # like LM_CpAIC_minus_corr

##
## Call:
## lm(formula = price ~ cut + color + clarity + carat + depth_ratio,
##     data = Train_minus_corr)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.98668 -0.08451 -0.00056  0.08207  1.42904
##
## Coefficients:
##              Estimate Std. Error t value      Pr(>|t|)
## (Intercept) -0.0733824  0.0015697  -46.750 < 0.0000000000000002 ***
## cut.L        0.1176217  0.0034546   34.048 < 0.0000000000000002 ***
## cut.Q       -0.0331066  0.0029605  -11.183 < 0.0000000000000002 ***
## cut.C        0.0133117  0.0025217    5.279   0.0000001310 ***
## cut^4       -0.0022101  0.0020180   -1.095     0.2734
## color.L      0.4322448  0.0028231 153.112 < 0.0000000000000002 ***
## color.Q     -0.0958088  0.0025845  -37.070 < 0.0000000000000002 ***
## color.C      0.0133162  0.0024237    5.494   0.0000000396 ***
## color^4      0.0125930  0.0022255    5.658   0.0000000154 ***
## color^5     -0.0001032  0.0021098   -0.049     0.9610
## color^6      0.0028362  0.0019159    1.480     0.1388
## clarity.L    0.9069041  0.0049997 181.393 < 0.0000000000000002 ***
## clarity.Q   -0.2447790  0.0046521  -52.616 < 0.0000000000000002 ***
## clarity.C    0.1346281  0.0039812   33.816 < 0.0000000000000002 ***
## clarity^4   -0.0686735  0.0031846  -21.565 < 0.0000000000000002 ***
## clarity^5    0.0288920  0.0025952   11.133 < 0.0000000000000002 ***
## clarity^6   -0.0024996  0.0022504   -1.111     0.2667
## clarity^7    0.0320720  0.0019873   16.138 < 0.0000000000000002 ***
## carat       1.0865705  0.0009207 1180.212 < 0.0000000000000002 ***
## depth_ratio -0.0016535  0.0008632   -1.916     0.0554 .
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.1313 on 26812 degrees of freedom
## Multiple R-squared:  0.9828, Adjusted R-squared:  0.9828
## F-statistic: 8.047e+04 on 19 and 26812 DF,  p-value: < 0.00000000000000022

GlobalCrit(LM_minus_corr)

##
## -----
## GLOBAL VARIABLE SELECTION PROCEDURE
##
## ( Data = Train_minus_corr )

```

```
##
## A = cut
## B = color
## C = clarity
## D = carat
## E = depth_ratio
## F = table
##
## Models      | Cp              | AIC              |
## -----|-----|-----|
## BCD          | 1686.62 ( 8) | 31157.06 ( 8) |
## ABCD         | - 7.32 ( 2) | 32800.58 ( 2) |
## ACDE         | 24920.99 (10) | 15162.85 ( 9) |
## BCDE         | 1514.54 ( 7) | 31319.38 ( 7) |
## BCDF         | 1498.62 ( 6) | 31334.46 ( 6) |
## ABCDE        | - 8.99 ( 1) | 32802.26 ( 1) |
## ABCDF        | - 5.79 ( 4) | 32799.06 ( 4) |
## ACDEF        | 24919.19 ( 9) | 15162.82 (10) |
## BCDEF        | 1158.93 ( 5) | 31658.16 ( 5) |
## ABCDEF       | - 7.00 ( 3) | 32800.27 ( 3) |
##
## -----
```

```
##
## -----
## GLOBAL VARIABLE SELECTION PROCEDURE
##
## ( Data = Train_minus_corr )
##
## A = cut
## B = color
## C = clarity
## D = carat
## E = depth_ratio
## F = table
##
## Models      | Cp              | AIC              |
## -----|-----|-----|
## BCD          | 1686.62 ( 8) | 31157.06 ( 8) |
## ABCD         | - 7.32 ( 2) | 32800.58 ( 2) |
## ACDE         | 24920.99 (10) | 15162.85 ( 9) |
## BCDE         | 1514.54 ( 7) | 31319.38 ( 7) |
## BCDF         | 1498.62 ( 6) | 31334.46 ( 6) |
## ABCDE        | - 8.99 ( 1) | 32802.26 ( 1) |
## ABCDF        | - 5.79 ( 4) | 32799.06 ( 4) |
## ACDEF        | 24919.19 ( 9) | 15162.82 (10) |
## BCDEF        | 1158.93 ( 5) | 31658.16 ( 5) |
## ABCDEF       | - 7.00 ( 3) | 32800.27 ( 3) |
##
## -----
```

```
LM_CpAIC_minus_corr = lm(price ~ . , data = Train_lr[, c(1:5, 9)])
summary(LM_CpAIC_minus_corr)
```

```
##
```

```
## Call:
## lm(formula = price ~ ., data = Train_lr[, c(1:5, 9)])
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.98668 -0.08451 -0.00056  0.08207  1.42904
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -0.0733824  0.0015697 -46.750 < 0.0000000000000002 ***
## cut.L         0.1176217  0.0034546  34.048 < 0.0000000000000002 ***
## cut.Q        -0.0331066  0.0029605 -11.183 < 0.0000000000000002 ***
## cut.C         0.0133117  0.0025217   5.279  0.0000001310 ***
## cut^4        -0.0022101  0.0020180  -1.095    0.2734
## color.L       0.4322448  0.0028231 153.112 < 0.0000000000000002 ***
## color.Q      -0.0958088  0.0025845 -37.070 < 0.0000000000000002 ***
## color.C       0.0133162  0.0024237   5.494  0.0000000396 ***
## color^4       0.0125930  0.0022255   5.658  0.0000000154 ***
## color^5      -0.0001032  0.0021098  -0.049    0.9610
## color^6       0.0028362  0.0019159   1.480    0.1388
## clarity.L     0.9069041  0.0049997 181.393 < 0.0000000000000002 ***
## clarity.Q    -0.2447790  0.0046521 -52.616 < 0.0000000000000002 ***
## clarity.C     0.1346281  0.0039812  33.816 < 0.0000000000000002 ***
## clarity^4    -0.0686735  0.0031846 -21.565 < 0.0000000000000002 ***
## clarity^5     0.0288920  0.0025952  11.133 < 0.0000000000000002 ***
## clarity^6    -0.0024996  0.0022504  -1.111    0.2667
## clarity^7     0.0320720  0.0019873  16.138 < 0.0000000000000002 ***
## carat        1.0865705  0.0009207 1180.212 < 0.0000000000000002 ***
## depth_ratio -0.0016535  0.0008632  -1.916    0.0554 .
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.1313 on 26812 degrees of freedom
## Multiple R-squared:  0.9828, Adjusted R-squared:  0.9828
## F-statistic: 8.047e+04 on 19 and 26812 DF, p-value: < 0.00000000000000022
```

```
kable_styling(
  kable(
    data.table(Model = c("LM_complete", "LM_forward_complete", "LM_backward_complete",
                        "LM_stepwise_complete", "LM_CpAIC_complete",
                        "LM_minus_corr", "LM_forward_minus_corr", "LM_backward_minus_corr",
                        "LM_stepwise_minus_corr", "LM_CpAIC_minus_corr"),
              Cut      = c("X", "X", "X", "X", "X", "X", "X", "X", "X", "X"),
              Color    = c("X", "X", "X", "X", "X", "X", "X", "X", "X", "X"),
              Clarity   = c("X", "X", "X", "X", "X", "X", "X", "X", "X", "X"),
              Carat     = c("X", "X", "X", "X", "X", "X", "X", "X", "X", "X"),
              Length    = c("X", "X", "X", "X", "X", " ", " ", " ", " ", " "),
              Width     = c("X", "X", " ", " ", " ", " ", " ", " ", " ", " "),
              Depth     = c("X", "X", "X", "X", "X", " ", " ", " ", " ", " "),
              `Depth Ratio` = c("X", "X", "X", "X", "X", "X", "X", "X", "X", "X"),
              Table     = c("X", "X", " ", " ", " ", " ", "X", "X", " ", " ")
    ),
  align = 'lcccccccc') %>% kable_classic(),
full_width = TRUE)
```

| Model                  | Cut | Color | Clarity | Carat | Length | Width | Depth | Depth<br>Ratio | Table |
|------------------------|-----|-------|---------|-------|--------|-------|-------|----------------|-------|
| LM_complete            | X   | X     | X       | X     | X      | X     | X     | X              | X     |
| LM_forward_complete    | X   | X     | X       | X     | X      | X     | X     | X              | X     |
| LM_backward_complete   | X   | X     | X       | X     | X      | X     | X     | X              |       |
| LM_stepwise_complete   | X   | X     | X       | X     | X      | X     | X     | X              |       |
| LM_CpAIC_complete      | X   | X     | X       | X     | X      | X     | X     | X              |       |
| LM_minus_corr          | X   | X     | X       | X     |        |       |       | X              | X     |
| LM_forward_minus_corr  | X   | X     | X       | X     |        |       |       | X              | X     |
| LM_backward_minus_corr | X   | X     | X       | X     |        |       |       | X              |       |
| LM_stepwise_minus_corr | X   | X     | X       | X     |        |       |       | X              |       |
| LM_CpAIC_minus_corr    | X   | X     | X       | X     |        |       |       | X              |       |

In both cases, the forward selection doesn't discard any variables, whereas backward, stepwise and global selections all choose the same model with less variables than initially.

Thus, we have four different models emerging. We will keep LM\_complete, LM\_CpAIC\_complete, LM\_minus\_corr and LM\_CpAIC\_minus\_corr and remove the other models which are duplicates.

```
rm(LM_forward_complete, LM_backward_complete, LM_stepwise_complete, LM_forward_minus_corr, LM_backward_minus_corr)
```

```
# predicting prices of validation set on the validation data
```

```
LM_Predictions =  
  data.table(  
    LM_complete_pred = predict(object = LM_complete, newdata = Valid_lr),  
    LM_CpAIC_complete_pred = predict(object = LM_CpAIC_complete, newdata = Valid_lr),  
    LM_minus_corr_pred = predict(object = LM_minus_corr, newdata = Valid_lr),  
    LM_CpAIC_minus_corr_pred = predict(object = LM_CpAIC_minus_corr, newdata = Valid_lr)  
  )
```

```
# we have to scale back the price, to do so we fetch the mean and std value from norm.values  
# displaying all means and stds  
norm.values$mean
```

```
##      price      carat      length      width      depth depth_ratio      table  
##  7.7876577 -0.3941013  5.7331235  5.7353183  3.5414058  61.7640019  4.0504070
```

```
norm.values$std
```

```
##      price      carat      length      width      depth depth_ratio      table  
##  1.01430151  0.58510643  1.12051551  1.11223832  0.69253953  1.41492652  0.03808011
```

```
# fetching for price
```

```
mean_price = norm.values$mean[1]  
std_price = norm.values$std[1]
```

```
# scaling back (Y*mu + sigma), then exp() (we had transformed price with a log for skewness)
```

```
LM_Predictions = LM_Predictions*std_price + mean_price  
LM_Predictions = exp(LM_Predictions)
```

```
# taking real prices of validation data from diamonds (which has not been touched -> original scale)  
LM_Predictions[, real_prices := diamonds[valid.index, price]]
```

```
Acc1 = accuracy(object = LM_Predictions$LM_complete_pred, x = LM_Predictions$real_prices)
```

| Model               | ME       | RMSE     | MAE      | MPE        | MAPE     |
|---------------------|----------|----------|----------|------------|----------|
| LM_complete         | 36.36648 | 846.5888 | 408.5364 | -0.8377036 | 10.33643 |
| LM_CpAIC_complete   | 36.32774 | 845.5018 | 408.1349 | -0.8408540 | 10.33659 |
| LM_minus_corr       | 50.38994 | 810.1522 | 405.0486 | -0.8420621 | 10.39559 |
| LM_CpAIC_minus_corr | 50.39878 | 810.1610 | 405.0616 | -0.8418121 | 10.39568 |

```

Acc1 = as.data.table(Acc1)
Acc2 = accuracy(object = LM_Predictions$LM_CpAIC_complete_pred, x = LM_Predictions$real_prices)
Acc2 = as.data.table(Acc2)
Acc3 = accuracy(object = LM_Predictions$LM_minus_corr_pred, x = LM_Predictions$real_prices)
Acc3 = as.data.table(Acc3)
Acc4 = accuracy(object = LM_Predictions$LM_CpAIC_minus_corr_pred, x = LM_Predictions$real_prices)
Acc4 = as.data.table(Acc4)
Accs = list(Acc1, Acc2, Acc3, Acc4)
Accs = rbindlist(Accs)
#rm(Acc1, Acc2, Acc3, Acc4)
Accs[, Model := c("LM_complete", "LM_CpAIC_complete", "LM_minus_corr", "LM_CpAIC_minus_corr")]
Accs <- Accs[, c(6, 1:5)]
kable_styling(kable(Accs)%>% kable_classic(), full_width = TRUE)

```

## $k$ -NN

```

diamonds_dummies <- diamonds
diamonds_dummies <- dummy_cols(diamonds_dummies,
                               select_columns = c("cut", "color", "clarity"),
                               remove_selected_columns = TRUE)

#rename column to avoid issues with Neural net function
colnames(diamonds_dummies)[10] = "cut_Very_Good"

#create data frames to normalize
knn_train <- diamonds_dummies[train.index,]
knn_valid <- diamonds_dummies[valid.index,]

knn_train_norm <- knn_train
knn_valid_norm <- knn_valid

# use preProcess() to normalize non-categorical variables
norm_values <- preProcess(knn_train[, c(1:7)], method=c("center", "scale"))

knn_train_norm[, c(1:7)] <- predict(norm_values, knn_train[, c(1:7)])
knn_valid_norm[,c(1:7)] <- predict(norm_values, knn_valid[, c(1:7)])

# computing kNN with knnreg from caret package
kNN = knnreg(x = knn_train_norm[, -c(1)] , y = knn_train_norm$price, k = 1)

# predicting prices of validation set on the validation data
knn_pred_y = predict(object = kNN, newdata = knn_valid_norm[, -c(1)])

# we have to scale back the price, to do so we fetch the mean and std value from norm.values
# displaying all means and standard deviations
norm_values$mean

```



|          | ME       | RMSE     | MAE      | MPE       | MAPE     |
|----------|----------|----------|----------|-----------|----------|
| Test set | 20.59665 | 954.5954 | 487.1664 | -1.833533 | 13.73317 |

```
##      price      carat      length      width      depth  depth_ratio      table
## 3934.4128280    0.7987388    5.7331235    5.7353183    3.5414058    61.7640019    57.4628205
norm_values$std
```

```
##      price      carat      length      width      depth  depth_ratio      table
## 3991.7064934    0.4742740    1.1205155    1.1122383    0.6925395    1.4149265    2.2176345
```

```
# fetching for price
mean_price = norm_values$mean[1]
std_price = norm_values$std[1]

# scaling back (Y*mu + sigma)
knn_rescaled_prices = std_price * knn_pred_y + mean_price

# taking real prices of validation data from diamonds (which has not been touched -> original scale)
real_prices = diamonds[valid.index, price]

# computing accuracy measures
kNN_aa1 <- accuracy(object = knn_rescaled_prices, x = real_prices)

kable_styling(kable(kNN_aa1)%>% kable_classic(), full_width = TRUE)

#create a data frame to store accuracy results
accuracy_df <- data.frame(k = seq(1, 20, 1), accuracy = rep(0, 20))

# use validation data set to compute various values of Knn
for(i in 1:20) {
  temp_knn_pred <- knnreg(x = knn_train_norm[, -c(1)] , y = knn_train_norm$price, k = i)

  temp_knn_pred_y = predict(object = temp_knn_pred, newdata = knn_valid_norm[, -c(1)])
  temp_knn_rescaled_prices = std_price * temp_knn_pred_y + mean_price

  #update the accuracy table with results of confusion matrix accuracy from the loop
  accuracy_df[i, 2] <- accuracy(object = temp_knn_rescaled_prices, x = real_prices)[2]
}

#display results
row_spec(kable_classic(kbl(accuracy_df)), 4, bold = T, color = "white", background = "#D8B365")

#plot results
ggplot (data = accuracy_df, aes(x = k, y = accuracy)) +
  geom_line (size = 1.2, color = "black") +
  geom_point(data= accuracy_df[4,], aes(x = k, y = accuracy), color = "#D8B365", size = 3) +
  labs(x = "Amount of Neighbors", y = "RMSE per Model", title = "Behavior of RMSE per amount of kNN")+
  theme_classic()

# computing kNN with knnreg from caret package
opt_kNN = knnreg(x = knn_train_norm[, -c(1)] , y = knn_train_norm$price, k = 4)

# predicting prices of validation set on the validation data
```

| k  | accuracy |
|----|----------|
| 1  | 954.5954 |
| 2  | 861.0887 |
| 3  | 827.1196 |
| 4  | 814.5940 |
| 5  | 819.2859 |
| 6  | 818.5541 |
| 7  | 822.7554 |
| 8  | 833.2104 |
| 9  | 839.9301 |
| 10 | 847.3985 |
| 11 | 850.7010 |
| 12 | 854.8280 |
| 13 | 860.4922 |
| 14 | 867.3541 |
| 15 | 874.8586 |
| 16 | 879.9069 |
| 17 | 886.0747 |
| 18 | 891.9780 |
| 19 | 896.3521 |
| 20 | 900.0938 |

|          | ME       | RMSE    | MAE      | MPE      | MAPE     |
|----------|----------|---------|----------|----------|----------|
| Test set | 40.86142 | 814.594 | 431.2913 | -2.45264 | 12.13511 |

```

opt_knn_pred_y = predict(object = opt_kNN, newdata = knn_valid_norm[, -c(1)])

# scaling back (Y*mu + sigma)
opt_knn_rescaled_prices = std_price * opt_knn_pred_y + mean_price

# computing accuracy measures
opt_accuracy <- accuracy(object = opt_knn_rescaled_prices, x = real_prices)
kable_styling(kable(opt_accuracy)%>% kable_classic(), full_width = TRUE)

```

## Regression Tree

### Partitioning

```

#Rename data specifically for regression trees
diamonds_tree <- diamonds

# Creating data tables Train, Valid and Test using the indexes for the regression tree section
Train_rg <- diamonds[train.index, ]
Valid_rg <- diamonds[valid.index, ]

```

### Regression Tree

```

# Use rpart() to run tree on continuous response
RegressTree <- rpart(price ~ .,
                     data = Train_rg,
                     method = "anova")

```

|          | ME       | RMSE     | MAE      | MPE       | MAPE     |
|----------|----------|----------|----------|-----------|----------|
| Test set | 6.097627 | 1267.154 | 846.9286 | -14.54158 | 33.08388 |

```
# Generates a cost complexity parameter table that provides the complexity parameter value
#summary(RegressTree)
```

```
# Plots a regression tree
```

```
fancyRpartPlot(RegressTree, caption = NULL, main = "Regression Tree", palettes = "YlGnBu", digits = -3)
```

```
# Count number of leaves
```

```
length(RegressTree$frame$var[RegressTree$frame$var == "<leaf>"])
```

```
## [1] 8
```

```
# kable and kable_styling as before
```

```
# We multiply by 100, divide by the sum and round the percentages to 2 decimals
```

```
kable_styling(kable(round(100*RegressTree$variable.importance / sum(RegressTree$variable.importance), 2),
```

|             | Importance % |
|-------------|--------------|
| width       | 25.50        |
| length      | 24.30        |
| carat       | 23.90        |
| depth       | 22.57        |
| clarity     | 2.53         |
| color       | 1.08         |
| depth_ratio | 0.09         |
| table       | 0.02         |
| cut         | 0.01         |

```
# Predict errors using accuracy()
```

```
tree_aa1 <- forecast::accuracy(predict(RegressTree, Valid_rg), Valid_rg$price)
```

```
kable_styling(kable(tree_aa1)%>% kable_classic(), full_width = TRUE)
```

```
# Predict the diamond price with validation
```

```
pred_Diamond_test <- predict(RegressTree, newdata = Valid_rg)
```

```
# Display first 14 observations
```

```
head(pred_Diamond_test,14)
```

```
##      1      2      3      4      5      6      7      8      9     10     11
## 5440.419 1059.257 1059.257 8560.343 1059.257 1059.257 3203.633 8414.236 3203.633 1059.257 5440.419 1
##      13      14
## 1059.257 5440.419
```

### ‘Exlcusion’ Regression Tree

```
RegressTree2 <- rpart(price ~ length+width+depth+carat+cut+color+clarity,
  data = Train_rg,
  method = "anova")
```

```
# Generates a cost complexity parameter table that provides the complexity parameter value
#summary(RegressTree2)
```

|          | ME       | RMSE     | MAE      | MPE       | MAPE     |
|----------|----------|----------|----------|-----------|----------|
| Test set | 6.097627 | 1267.154 | 846.9286 | -14.54158 | 33.08388 |

```
# Plots a regression tree
fancyRpartPlot(RegressTree2, caption = NULL, main = "Exclusion Regression Tree", palettes = "YlGnBu", d
```

```
# kable and kable_styling as before
# We multiply by 100, divide by the sum and round the percentages to 2 decimals
kable_styling(kable(round(100*RegressTree2$variable.importance / sum(RegressTree2$variable.importance),
```

|         | Importance % |
|---------|--------------|
| width   | 25.51        |
| length  | 24.31        |
| carat   | 23.91        |
| depth   | 22.58        |
| clarity | 2.55         |
| color   | 1.08         |
| cut     | 0.05         |

```
# Predict errors using accuracy()
tree_aa2 <- forecast::accuracy(predict(RegressTree2, Valid_rg), Valid_rg$price)

kable_styling(kable(tree_aa2)%>% kable_classic(), full_width = TRUE)
```

```
# Predict the diamond price with validation
pred_Diamond_test <- predict(RegressTree, newdata = Valid_rg)
```

```
# Display first 14 observations
head(pred_Diamond_test,14)
```

```
##          1          2          3          4          5          6          7          8          9         10         11
## 5440.419 1059.257 1059.257 8560.343 1059.257 1059.257 3203.633 8414.236 3203.633 1059.257 5440.419 1
##          13          14
## 1059.257 5440.419
```

```
#Get the lowest CP value from CP table
min.xerror <- RegressTree2$cptable[which.min(RegressTree2$cptable[, "xerror"]), "CP"]
```

```
min.xerror
```

```
## [1] 0.01
```

```
#Plot the optimal Cp value
plotcp(RegressTree2)
```

## Pruned Regression Tree

```
RegressTree_pruned <- prune(RegressTree2, cp = min.xerror)
```

```
# Draw the prune tree
fancyRpartPlot(RegressTree_pruned, caption = NULL, main = "'Pruned' Regression Tree", palettes = "YlGnBu"
```

|          | ME       | RMSE     | MAE      | MPE       | MAPE     |
|----------|----------|----------|----------|-----------|----------|
| Test set | 6.097627 | 1267.154 | 846.9286 | -14.54158 | 33.08388 |

|          | ME        | RMSE     | MAE      | MPE       | MAPE     |
|----------|-----------|----------|----------|-----------|----------|
| Test set | -3.898022 | 3660.422 | 2765.963 | -144.4894 | 171.6511 |

```
# Predict errors using accuracy()
tree_aa3 <- forecast::accuracy(predict(RegressTree_pruned, Valid_rg), Valid_rg$price)

kable_styling(kable(tree_aa3)%>% kable_classic(), full_width = TRUE)
```

## Boosted Tree

```
# Boosted tree
set.seed(111)
tree_boost10 <- gbm(price ~., data = Train_rg, distribution = "gaussian", n.trees = 10, interaction.depth = 6, shrinkage = 0.01)

# Interaction Depth specifies the maximum depth of each tree( i.e. highest level of variable interaction)
# Shrinkage is considered as the learning rate. It is used for reducing, or shrinking, the impact of each variable.
# n.trees: Integer specifying the total number of trees to fit. This is equivalent to the number of iterations.

tree_boost10

## gbm(formula = price ~ ., distribution = "gaussian", data = Train_rg,
##      n.trees = 10, interaction.depth = 6, shrinkage = 0.01)
## A gradient boosted model with gaussian loss function.
## 10 iterations were performed.
## There were 9 predictors of which 5 had non-zero influence.

vip::vip(tree_boost10, aesthetics = list(fill = "#D8B365")) +
  ggtitle("Boosted Tree Variable Importance") + # Title name
  xlab("Variable") + # Label names
  theme_classic() # A classic theme, with x and y axis lines and no grid lines
```

```
# Predict errors using accuracy()
tree_aa4 <- forecast::accuracy(predict(tree_boost10, Valid_rg), Valid_rg$price)

kable_styling(kable(tree_aa4)%>% kable_classic(), full_width = TRUE)
```

```
# Boosted tree
set.seed(111)
tree_boost30 <- gbm(price ~., data = Train_rg, distribution = "gaussian", n.trees = 30)

vip::vip(tree_boost30, aesthetics = list(fill = "#D8B365")) +
  ggtitle("Boosted Tree Variable Importance") + # Title name
  xlab("Variable") + # Label names
  theme_classic() # A classic theme, with x and y axis lines and no grid lines
```

|          | ME     | RMSE     | MAE      | MPE       | MAPE     |
|----------|--------|----------|----------|-----------|----------|
| Test set | 2.5438 | 1531.521 | 985.6974 | -35.25731 | 46.85606 |

|          | ME       | RMSE     | MAE      | MPE       | MAPE     |
|----------|----------|----------|----------|-----------|----------|
| Test set | 9.002539 | 1170.295 | 680.2182 | -12.85125 | 26.22023 |

```
tree_boost30
```

```
## gbm(formula = price ~ ., distribution = "gaussian", data = Train_rg,
##     n.trees = 30)
## A gradient boosted model with gaussian loss function.
## 30 iterations were performed.
## There were 9 predictors of which 4 had non-zero influence.
```

```
# Predict errors using accuracy()
tree_aa5 <- forecast::accuracy(predict(tree_boost30, Train_rg), Train_rg$price)
```

```
kable_styling(kable(tree_aa5)%>% kable_classic(), full_width = TRUE)
```

```
# Boosted tree
```

```
set.seed(111)
tree_boost100 <- gbm(price ~., data = Train_rg, distribution = "gaussian", cv.folds = 3)
```

```
# cv.folds: Number of cross-validation folds to perform. If cv.folds>1 then gbm, in addition to the usual
```

```
vip::vip(tree_boost100, aesthetics = list(fill = "#D8B365")) +
  ggtitle("Boosted Tree Variable Importance") + # Title name
  xlab("Variable") + # Label names
  theme_classic() # A classic theme, with x and y axis lines and no grid lines
```

```
tree_boost100
```

```
## gbm(formula = price ~ ., distribution = "gaussian", data = Train_rg,
##     cv.folds = 3)
## A gradient boosted model with gaussian loss function.
## 100 iterations were performed.
## The best cross-validation iteration was 100.
## There were 9 predictors of which 6 had non-zero influence.
```

```
# Predict errors using accuracy()
tree_aa6 <- forecast::accuracy(predict(tree_boost100, Valid_rg), Valid_rg$price)
```

```
kable_styling(kable(tree_aa6)%>% kable_classic(), full_width = TRUE)
```

```
pred_boost <- predict.gbm(tree_boost100, newdata = Valid_rg)
```

```
head(pred_boost, 14)
```

```
## [1] 6112.0846 1582.6043 743.7071 6352.0316 1140.2701 1491.0859 2806.4714 9855.1238 2681.4129 1330.4
## [11] 7595.4454 1741.6017 1380.0364 6132.5267
```

|          | ME        | RMSE     | MAE      | MPE       | MAPE     |
|----------|-----------|----------|----------|-----------|----------|
| Test set | 0.9333074 | 1248.769 | 808.9297 | -14.35086 | 31.45019 |

## Bagging Tree

```
# Bagged tree
set.seed(111)
tree_bagging <- bagging(price ~., data = Train_rg, coob = TRUE)
tree_bagging

##
## Bagging regression trees with 25 bootstrap replications
##
## Call: bagging.data.frame(formula = price ~ ., data = Train_rg, coob = TRUE)
##
## Out-of-bag estimate of root mean squared error: 1268.231
ss<-varImp(tree_bagging)

data<-data.table(name=row.names(ss),value=ss$Overall)

data[,ggplot(.SD, aes(x=reorder(name, ss$Overall), y=ss$Overall)) +
  geom_bar(stat = "identity", fill = "#D8B365") +
  xlab("Variable") + # Label names +
  ylab("Importance") + # Label names +
  ggtitle("Bagged Tree Variable Importance") + # Title name
  theme_classic() + # A classic theme, with x and y axis lines and no grid lines
  coord_flip(),]
```

```
tree_aa7 <- forecast::accuracy(predict(tree_bagging, Valid_rg), Valid_rg$price)

kable_styling(kable(tree_aa7)%>% kable_classic(), full_width = TRUE)
```

```
pred_bagged <- predict(tree_bagging, newdata = Valid_rg)

head(pred_bagged,14)
```

```
## [1] 5396.881 1104.498 1046.993 8084.493 1104.498 1104.498 3123.791 10274.884 3123.791 1046
## [11] 6052.800 1104.498 1046.993 5953.609
```

## RandomForest

```
options(scipen = 9999)

# Create randomForest
set.seed(111)
rdf_model <- randomForest(price~ ., ntree= 60, data = Train_rg)
rdf_model

##
## Call:
## randomForest(formula = price ~ ., data = Train_rg, ntree = 60)
##
## Type of random forest: regression
```

|          | ME       | RMSE   | MAE      | MPE       | MAPE     |
|----------|----------|--------|----------|-----------|----------|
| Test set | 2.969775 | 577.63 | 289.5139 | -1.402886 | 7.235286 |

```
##                      Number of trees: 60
## No. of variables tried at each split: 3
##
##          Mean of squared residuals: 335950
##                      % Var explained: 97.89
```

```
plot(rdf_model,
      main = "Number of Optimal Trees")
```

```
# Number of trees with lowest MSE
# Compared with 80 and 100 trees, the trade-off with both of those were small so we stuck with 60 trees
which.min(rdf_model$mse)
```

```
## [1] 59
```

```
options(scipen = 9999)
```

```
# Create variable importance chart
vip::vip(rdf_model, aesthetics = list(fill = "#D8B365")) +
  ggtitle("RandomTree Variable Importance") + # Title name
  xlab("Variable") + # Label names
  theme_classic() # A classic theme, with x and y axis lines and no grid lines
```

```
tree_aa8 <- forecast::accuracy(predict(rdf_model, Valid_rg), Valid_rg$price)
```

```
kable_styling(kable(tree_aa8)%>% kable_classic(), full_width = TRUE)
```

```
pred_random <- predict(rdf_model, newdata = Valid_rg)
```

```
head(pred_random,14)
```

```
##          1          2          3          4          5          6          7          8          9         10
## 5012.2100 1782.6600  591.9428 9185.5925 1319.5133 1344.4381 2423.8250 8620.2650 2562.5817 936.9519
##          12          13          14
## 1788.5992 1172.3436 7014.9981
```

## Create Regression Tree Summary Table

```
Tree_res <- data.frame("Model_Tree" = c("Reg. Tree",
                                         "Exclus. RT",
                                         "Pruned RT",
                                         "Boost10",
                                         "Boost30",
                                         "Boost100",
                                         "Bagging",
                                         "RndmFrst"),
                      "ME" = "",
                      "RMSE" = "",
                      "MAE" = "",
```



|            | ME   | RMSE    | MAE     | MPE     | MAPE   |
|------------|------|---------|---------|---------|--------|
| Reg. Tree  | 6.1  | 1267.15 | 846.93  | -14.54  | 33.08  |
| Exclus. RT | 6.1  | 1267.15 | 846.93  | -14.54  | 33.08  |
| Pruned RT  | 6.1  | 1267.15 | 846.93  | -14.54  | 33.08  |
| Boost10    | -3.9 | 3660.42 | 2765.96 | -144.49 | 171.65 |
| Boost30    | 2.54 | 1531.52 | 985.7   | -35.26  | 46.86  |
| Boost100   | 9    | 1170.3  | 680.22  | -12.85  | 26.22  |
| Bagging    | 0.93 | 1248.77 | 808.93  | -14.35  | 31.45  |
| RndmFrst   | 2.97 | 577.63  | 289.51  | -1.4    | 7.24   |

```

      "MPE" = "",
      "MAPE" = "")

rownames(Tree_res) <- Tree_res$Model_Tree
Tree_res$Model_Tree = NULL

for (i in 1:8) {
  Tree_res[i,] = round(get(paste("tree_aa", i, sep = ""), 2), 2)
}

kable_styling(kable(Tree_res)%>% kable_classic(), full_width = TRUE)

RGRMSEplotdata <- t(Tree_res$RMSE)
colnames(RGRMSEplotdata) <- t(sort(rownames(Tree_res), decreasing = TRUE))
RGRMSEplotdata

##      RndmFrst  Reg. Tree Pruned RT Exclus. RT Boost30  Boost100 Boost10  Bagging
## [1,] "1267.15" "1267.15" "1267.15" "3660.42"  "1531.52" "1170.3" "1248.77" "577.63"

barplot(RGRMSEplotdata, col = "#D8B365", border = "#D8B365" ,
        main = "Regression Tree Comparison", ylab = "RMSE", las = 3, ylim=c(0,4000))

```

## NeuralNetworks

### Data preprocessing

```

#create dummy columns
diamonds_dummies <- dummy_cols(diamonds, select_columns = c("cut", "color", "clarity"))
diamonds_dummies <- diamonds_dummies[, -c(2:4)]

#rename column to avoid issues with Neural net function
colnames(diamonds_dummies)[10] = "cut_Very_Good"

```

### Partitioning

```

Train_nn <- diamonds_dummies[train.index,]
Valid_nn <- diamonds_dummies[valid.index,]

```

## Normalizing Data

```
# use preProcess() to normalize non-categorical variables
norm_values <- preProcess(Train_nn[1:7], method= c("range"))

Train_nn_norm <- predict(norm_values, Train_nn)
Valid_nn_norm <- predict(norm_values, Valid_nn)
```

## Model 1 (Layers = 1, Nodes = 1)

- Creating Model 1

```
#create data sets to train and validate the model
#train
x_train <- c(t(Train_nn_norm[, -c(1)]))
x_train <- as.array(x_train,
                    dim(t(Train_nn_norm[, -c(1)])),
                    dimnames = list(rownames(x_train), colnames(x_train)))
x_train <- as_tensor(x_train, shape = dim(Train_nn_norm[, -c(1)]))
y_train <- as_tensor(Train_nn_norm$price)

#validation
x_valid <- c(t(Valid_nn_norm[, -c(1)]))
x_valid <- as.array(x_valid,
                    dim(t(Valid_nn_norm[, -c(1)])),
                    dimnames = list(rownames(x_valid), colnames(x_valid)))
x_valid <- as_tensor(x_valid, shape = dim(Valid_nn_norm[, -c(1)]))
y_valid <- as_tensor(Valid_nn_norm$price)

#create model
rm(model_1_1) #to prevent retraining of existing model
tf$random$set_seed(111)
model_1_1 <- keras_model_sequential(input_shape = ncol(x_train)) %>%
  layer_dense(1, activation = "sigmoid", name = "HiddenLayer") %>% # 1st hidden layer (26 nodes)
  layer_dense(1, activation = "sigmoid", name = "outputLayer") # output layer (1 node)

#compile the model to be trained
model_1_1 %>% compile(
  optimizer = optimizer_adam(),
  loss = loss_mean_squared_error(),
  metric = metric_root_mean_squared_error()
)

#train the model and find the lowest rmse
history <- model_1_1 %>% fit(
  x = x_train,
  y = y_train,
  validation_data = c(x_valid, y_valid),
  verbose = FALSE
)
```

- Measuring Error Model 1

|          | ME        | RMSE     | MAE      | MPE       | MAPE     |
|----------|-----------|----------|----------|-----------|----------|
| Test set | -6.616364 | 3982.992 | 3025.611 | -157.4749 | 187.2922 |

```
# prediction values of validation set
y_valid_pred <- predict(model_1_1, x_valid)

# fetching a and b from standardization for scaling back price
a <- norm_values$ranges[1,1]
b <- norm_values$ranges[2,1]

valRescale <- function (x) {
  value <- (x * (b-a) + a)
  return(value)
}

# scaling back price predictions
y_valid_pred <- valRescale(y_valid_pred)

# we have to change the class of y_valid_pred
class(y_valid_pred)

## [1] "matrix" "array"

y_valid_pred <- as.numeric(y_valid_pred)
class(y_valid_pred)

## [1] "numeric"

# taking real values of validation set from original data (which wasn't standardized!)
y_valid_real <- diamonds[valid.index, price]

# computing accuracy measures of validation set
acc_1 <- accuracy(object = y_valid_pred, x = y_valid_real)

kable_styling(kable(acc_1)%>% kable_classic(), full_width = TRUE)
```

## Model 2 (Layers = 1, Nodes = 26)

- Creating Model 2

```
#create model
rm(model_1_26) #to prevent retraining of existing model
tf$random$set_seed(111)
model_1_26 <- keras_model_sequential(input_shape = ncol(x_train)) %>%
  layer_dense(26, activation = "sigmoid", name = "HiddenLayer") %>% # 1st hidden layer (26 nodes)
  layer_dense(1, activation = "sigmoid", name = "outputLayer") # output layer (1 node)

#compile the model to be trained
model_1_26 %>% compile(
  optimizer = optimizer_adam(),
  loss = loss_mean_squared_error(),
  metric = metric_root_mean_squared_error()
)
```

|          | ME        | RMSE     | MAE     | MPE       | MAPE     |
|----------|-----------|----------|---------|-----------|----------|
| Test set | -83.27564 | 745.1569 | 450.941 | -11.36538 | 17.91012 |

```
#train the model and find the lowest rmse
history <- model_1_26 %>% fit(
  x = x_train,
  y = y_train,
  validation_data = c(x_valid, y_valid),
  verbose = FALSE
)
```

- Measuring Error Model 2

```
# prediction values of validation set
y_valid_pred_1_26 <- predict(model_1_26, x_valid)

# scaling back price predictions
y_valid_pred_1_26 <- valRescale(y_valid_pred_1_26)

# we have to change the class of y_valid_pred
y_valid_pred_1_26 <- as.numeric(y_valid_pred_1_26)

# computing accuracy measures of validation set
acc_2 <- accuracy(object = y_valid_pred_1_26, x = y_valid_real)

kable_styling(kable(acc_2)%>% kable_classic(), full_width = TRUE)
```

### Model 3 (Layers = 1, Nodes = 13)

- Creating Model 3

```
#create model
rm(model_1_13) #to prevent retraining of existing model
tf$random$set_seed(111)
model_1_13 <- keras_model_sequential(input_shape = ncol(x_train)) %>%
  layer_dense(13, activation = "sigmoid", name = "HiddenLayer") %>% # 1st hidden layer (26 nodes)
  layer_dense(1, activation = "sigmoid", name = "outputLayer") # output layer (1 node)

#compile the model to be trained
model_1_13 %>% compile(
  optimizer = optimizer_adam(),
  loss = loss_mean_squared_error(),
  metric = metric_root_mean_squared_error()
)

#train the model and find the lowest rmse
history <- model_1_13 %>% fit(
  x = x_train,
  y = y_train,
  validation_data = c(x_valid, y_valid),
  verbose = FALSE
)
```

- Measuring Error Model 3

|          | ME        | RMSE     | MAE      | MPE      | MAPE     |
|----------|-----------|----------|----------|----------|----------|
| Test set | -86.71198 | 744.0357 | 454.3788 | -12.8878 | 19.21285 |

```
# prediction values of validation set
y_valid_pred_1_13 <- predict(model_1_13, x_valid)

# scaling back price predictions
y_valid_pred_1_13 <- valRescale(y_valid_pred_1_13)

# we have to change the class of y_valid_pred
y_valid_pred_1_13 <- as.numeric(y_valid_pred_1_13)

# computing accuracy measures of validation set
acc_3 <- accuracy(object = y_valid_pred_1_13, x = y_valid_real)

kable_styling(kable(acc_3)%>% kable_classic(), full_width = TRUE)
```

#### Model 4 (Layers = 2, Nodes = 26)

- Creating Model 4

```
#create model
rm(model_2_26) #to prevent retraining of existing model
tf$random$set_seed(111)
model_2_26 <- keras_model_sequential(input_shape = ncol(x_train)) %>%
  layer_dense(26, activation = "sigmoid", name = "HiddenLayer") %>% # 1st hidden layer (26 nodes)
  layer_dense(26, activation = "sigmoid", name = "HiddenLayer2") %>% # 2nd hidden layer (26 nodes)
  layer_dense(13, activation = "sigmoid", name = "HiddenLayer3") %>% # 3rd hidden layer (26 nodes)
  layer_dense(1, activation = "sigmoid", name = "outputLayer") # output layer (1 node)

#compile the model to be trained
model_2_26 %>% compile(
  optimizer = optimizer_adam(),
  loss = loss_mean_squared_error(),
  metric = metric_root_mean_squared_error()
)

#train the model and find the lowest rmse
history <- model_2_26 %>% fit(
  x = x_train,
  y = y_train,
  validation_data = c(x_valid, y_valid),
  verbose = FALSE
)
```

- Measuring Error Model 4

```
# prediction values of validation set
y_valid_pred_2_26 <- predict(model_2_26, x_valid)

# scaling back price predictions
y_valid_pred_2_26 <- valRescale(y_valid_pred_2_26)
```

|          | ME       | RMSE    | MAE      | MPE       | MAPE     |
|----------|----------|---------|----------|-----------|----------|
| Test set | 30.21931 | 627.652 | 364.5433 | -5.429278 | 13.48542 |

```
# we have to change the class of y_valid_pred
y_valid_pred_2_26 <- as.numeric(y_valid_pred_2_26)

# computing accuracy measures of validation set
acc_4 <- accuracy(object = y_valid_pred_2_26, x = y_valid_real)

kable_styling(kable(acc_4)%>% kable_classic(), full_width = TRUE)
```

### Model 5 with Initializer (Layers = 2, Nodes = 26, GlorotNormal)

- Creating Model 5

```
#create model
rm(model_2_26G) #to prevent retraining of existing model
tf$random$set_seed(111)
model_2_26G <- keras_model_sequential(input_shape = ncol(x_train)) %>%
  layer_dense(26, activation = "sigmoid", name = "HiddenLayer", kernel_initializer = "GlorotNormal") %>%
  layer_dense(26, activation = "sigmoid", name = "HiddenLayer2", kernel_initializer = "GlorotNormal") %>%
  layer_dense(26, activation = "sigmoid", name = "HiddenLayer3", kernel_initializer = "GlorotNormal") %>%
  layer_dense(1, activation = "sigmoid", name = "outputLayer") # output layer (1 node)

#compile the model to be trained
model_2_26G %>% compile(
  optimizer = optimizer_adam(),
  loss = loss_mean_squared_error(),
  metric = metric_root_mean_squared_error()
)

#train the model and find the lowest rmse
history <- model_2_26G %>% fit(
  x = x_train,
  y = y_train,
  validation_data = c(x_valid, y_valid),
  verbose = FALSE
)
```

- Measuring Error Model 5

```
# prediction values of validation set
y_valid_pred_2_26G <- predict(model_2_26G, x_valid)

# scaling back price predictions
y_valid_pred_2_26G <- valRescale(y_valid_pred_2_26G)

# we have to change the class of y_valid_pred
y_valid_pred_2_26G <- as.numeric(y_valid_pred_2_26G)

# computing accuracy measures of validation set
acc_5 <- accuracy(object = y_valid_pred_2_26G, x = y_valid_real)
```

|          | ME        | RMSE     | MAE      | MPE       | MAPE     |
|----------|-----------|----------|----------|-----------|----------|
| Test set | -32.61097 | 630.0157 | 358.2125 | -6.623257 | 12.71689 |

|          | ME       | RMSE     | MAE      | MPE      | MAPE     |
|----------|----------|----------|----------|----------|----------|
| Test set | 8.980655 | 577.0847 | 322.9464 | -2.03136 | 10.22919 |

```
kable_styling(kable(acc_5)%>% kable_classic(), full_width = TRUE)
```

**Model 6 with Initializer and Learning Rate (Layers = 2, Nodes = 26, GlorotNormal, LR = 0.005)**

- Creating model 6

```
#create model
rm(model_2_26GLR) #to prevent retraining of existing model
tf$random$set_seed(111)
model_2_26GLR <- keras_model_sequential(input_shape = ncol(x_train)) %>%
  layer_dense(26, activation = "sigmoid", name = "HiddenLayer", kernel_initializer = "GlorotNormal") %>%
  layer_dense(26, activation = "sigmoid", name = "HiddenLayer2", kernel_initializer = "GlorotNormal") %>%
  layer_dense(13, activation = "sigmoid", name = "HiddenLayer3", kernel_initializer = "GlorotNormal") %>%
  layer_dense(1, activation = "sigmoid", name = "outputLayer")      # output layer (1 node)

#compile the model to be trained
model_2_26GLR %>% compile(
  optimizer = optimizer_adam(learning_rate = 0.009),
  loss = loss_mean_squared_error(),
  metric = metric_root_mean_squared_error()
)

#train the model and find the lowest rmse
history <- model_2_26GLR %>% fit(
  x = x_train,
  y = y_train,
  validation_data = c(x_valid, y_valid),
  verbose = FALSE
)
```

- Measuring Error Model 6

```
# prediction values of validation set
y_valid_pred_2_26GLR <- predict(model_2_26GLR, x_valid)

# scaling back price predictions
y_valid_pred_2_26GLR <- valRescale(y_valid_pred_2_26GLR)

# we have to change the class of y_valid_pred
y_valid_pred_2_26GLR <- as.numeric(y_valid_pred_2_26GLR)

# computing accuracy measures of validation set
acc_6 <- accuracy(object = y_valid_pred_2_26GLR, x = y_valid_real)

kable_styling(kable(acc_6)%>% kable_classic(), full_width = TRUE)
```

|             | ME     | RMSE    | MAE     | MPE     | MAPE   |
|-------------|--------|---------|---------|---------|--------|
| L1 N1       | -6.62  | 3982.99 | 3025.61 | -157.47 | 187.29 |
| L1 N26      | -83.28 | 745.16  | 450.94  | -11.37  | 17.91  |
| L1 N13      | -86.71 | 744.04  | 454.38  | -12.89  | 19.21  |
| L2 N26      | 30.22  | 627.65  | 364.54  | -5.43   | 13.49  |
| L2 N26 G    | -32.61 | 630.02  | 358.21  | -6.62   | 12.72  |
| L2 N26 G LR | 8.98   | 577.08  | 322.95  | -2.03   | 10.23  |

## Create Neural Net Summary Table

```

NN_res <- data.frame("Model" = c("L1 N1",
                                "L1 N26",
                                "L1 N13",
                                "L2 N26",
                                "L2 N26 G",
                                "L2 N26 G LR"),
                    "ME" = "",
                    "RMSE" = "",
                    "MAE" = "",
                    "MPE" = "",
                    "MAPE" = "")

rownames(NN_res) <- NN_res$Model
NN_res$Model = NULL

for (i in 1:6) {
  NN_res[i,] = round(get(paste("acc_", i, sep = "")), 2)
}

kable_styling(kable(NN_res)%>% kable_classic(), full_width = FALSE)

NNRMSEplotdata <- t(NN_res$RMSE)
colnames(NNRMSEplotdata) <- t(rownames(NN_res))

barplot(NNRMSEplotdata, col = "#D8B365", border = "#D8B365",
        main = "Neural Network Comparison", ylab = "RMSE", las = 3, ylim=c(0,4000))

```

## Neural network graphic description

### Ensembles

```

#create Table with 0 values
ensemble_summ <- data.frame("MLR" = 0,
                            "RegTrees" = 0,
                            "kNN" = 0,
                            "NeuralNets" = 0)

#create empty rows
ensemble_summ[length(valid.index),] <- 0

#Add values per columns
ensemble_summ[, "MLR"] = LM_Predictions$LM_CpAIC_minus_corr_pred
ensemble_summ[, "RegTrees"] = pred_random

```






Figure 1: Visual description of Neural Network

|          | ME       | RMSE     | MAE      | MPE       | MAPE     |
|----------|----------|----------|----------|-----------|----------|
| Test set | 25.80266 | 582.3353 | 306.1574 | -1.682175 | 8.448477 |

```
ensemble_summ[, "kNN"] = opt_knn_rescaled_prices
ensemble_summ[, "NeuralNets"] = y_valid_pred_2_26GLR

#Calculate Average
ensemble_summ$AveragePred <- rowMeans(ensemble_summ[, 1:4])
ensemble_summ$RealPrices <- real_prices

ens_error = accuracy(ensemble_summ$AveragePred, ensemble_summ$RealPrices)

kable_styling(kable(ens_error)%>% kable_classic(), full_width = TRUE)
```

```
ens_rmse_summ <- NN_res[0,] #NeuralNet Error
ens_rmse_summ[1, ] <- round(Acc4, 2)
ens_rmse_summ[2, ] <- Tree_res[8,]
ens_rmse_summ[3, ] <- round(opt_accuracy, 2)
ens_rmse_summ[4, ] <- NN_res[6,]
ens_rmse_summ[5, ] <- round(ens_error, 2)

row.names(ens_rmse_summ) <- c("Multiple Linear Regression",
                             "Regression Tree",
                             "k-Nearest Neighbor",
                             "Neural Network",
                             "Ensemble")

#display results
kable_styling(kable(ens_rmse_summ)%>% kable_classic(), full_width = FALSE)
```

## Model Performance Summary

```
predtable_lr <- data.frame("Real" = real_prices,
                          "Predicted" = ensemble_summ[, 1],
                          "Color" = diamonds[valid.index, 3],
                          "Cut" = diamonds[valid.index, 2],
                          "Clarity" = diamonds[valid.index, 4])
predtable_rg <- data.frame("Real" = real_prices,
                          "Predicted" = ensemble_summ[, 2],
                          "Color" = diamonds[valid.index, 3],
                          "Cut" = diamonds[valid.index, 2],
                          "Clarity" = diamonds[valid.index, 4])
predtable_knn <- data.frame("Real" = real_prices,
                          "Predicted" = ensemble_summ[, 3],
                          "Color" = diamonds[valid.index, 3],
                          "Cut" = diamonds[valid.index, 2],
                          "Clarity" = diamonds[valid.index, 4])
predtable_nn <- data.frame("Real" = real_prices,
                          "Predicted" = ensemble_summ[, 4],
                          "Color" = diamonds[valid.index, 3],
                          "Cut" = diamonds[valid.index, 2],
                          "Clarity" = diamonds[valid.index, 4])
```

```

predtable_ens <- data.frame("Real" = real_prices,
                           "Predicted" = ensemble_summ[, 5],
                           "Color" = diamonds[valid.index, 3],
                           "Cut" = diamonds[valid.index, 2],
                           "Clarity" = diamonds[valid.index, 4])

pred_plot_lr <- ggplot(data = predtable_lr, aes(x= Real, y = Predicted, color = clarity))+
  geom_point(show.title = FALSE, size = 0.5)+
  scale_color_brewer(type = 'div', guide = guide_legend(reverse = T, override.aes = list(alpha = 1, size = 0.5))+
  labs(x= "Multiple Linear Regression", y = "Predicted Values")+
  theme(legend.key.size = unit(0.5, 'cm'))+
  theme_classic()

pred_plot_rg <- ggplot(data = predtable_rg, aes(x= Real, y = Predicted, color = clarity))+
  geom_point(size = 0.5)+
  scale_color_brewer(type = 'div', guide = guide_legend(reverse = T, override.aes = list(alpha = 1, size = 0.5))+
  labs(x= "NN (Regression Trees", y = "")+
  theme_classic()

pred_plot_knn <- ggplot(data = predtable_knn, aes(x= Real, y = Predicted, color = clarity))+
  geom_point(size = 0.5)+
  scale_color_brewer(type = 'div', guide = guide_legend(reverse = T, override.aes = list(alpha = 1, size = 0.5))+
  labs(x= "k Neares Neighbor", y = "")+
  theme_classic()

pred_plot_nn <- ggplot(data = predtable_nn, aes(x= Real, y = Predicted, color = clarity))+
  geom_point(size = 0.5)+
  scale_color_brewer(type = 'div', guide = guide_legend(reverse = T, override.aes = list(alpha = 1, size = 0.5))+
  labs(x= "Neural Networks", y = "")+
  theme_classic()

pred_plot_ens <- ggplot(data = predtable_ens, aes(x= Real, y = Predicted, color = clarity))+
  geom_point(size = 0.5)+
  scale_color_brewer(type = 'div', guide = guide_legend(reverse = T, override.aes = list(alpha = 1, size = 0.5))+
  labs(x= "Ensemble", y = "")+
  theme_classic()

#pred_plot_full <- grid.arrange(ggplotGrob("Hello"),ggarrange(pred_plot_lr, pred_plot_rg, pred_plot_knn, pred_plot_nn, pred_plot_ens))
#pred_plot_full

```

## Conclusions